

The toy example: A Worked Analysis

Thomas E. Love, Ph.D.

Version: 2019-01-06

Contents

Setup	3
The Data Set	3
The Codebook for the <code>toy</code> data	4
“Skimmed” Summaries, within treatment groups	4
Table 1	6
Data Management and Cleanup	6
Range Checks for Quantitative (continuous) Variables	6
Restating Categorical Information in Helpful Ways	6
Dealing with Variables including More than Two Categories	8
Creating the Transformation and Product Terms	9
Data Set After Cleaning	9
Skim, within Treatment Groups	9
Table 1	10
The 13 Tasks We’ll Tackle in this Example	11
Task 1. Ignoring covariates, estimate the effect of treatment vs. control on...	12
Outcome 1 (a continuous outcome)	12
Outcome 2 (a binary outcome)	13
Outcome 3 (a time-to-event outcome with right censoring)	16
Unadjusted Estimates of Treatment Effect on Outcomes	18
Task 2. Fit the propensity score model, then plot the PS-treatment relationship	18
Comparing the Distribution of Propensity Score Across the Two Treatment Groups	19
Task 3. Rubin’s Rules to Check Overlap Before Propensity Adjustment	22
Rubin’s Rule 1	23
Rubin’s Rule 2	23
Rubin’s Rule 3	23
Task 4. Use 1:1 greedy matching on the linear PS, then check post-match balance	25
Balance Assessment (Semi-Automated)	26
Extracting, Tabulating Standardized Differences (without <code>cobalt</code>)	30
A Love Plot describing Standardized Differences Before/After Matching (without <code>cobalt</code>)	31
Using <code>cobalt</code> to build a “Love Plot” after Matching	32
Extracting, Tabulating Variance Ratios (without <code>cobalt</code>)	34
Creating a New Data Frame, Containing the Matched Sample (without <code>cobalt</code>)	35
Rubin’s Rules to Check Balance After Matching	35
Task 5. After matching, estimate the causal effect of treatment on ...	38
Outcome 1 (a continuous outcome)	38
Outcome 2 (a binary outcome)	41
Outcome 3 (a time-to-event outcome)	42
Results So Far (After Propensity Matching)	44

Task 6. Subclassify by PS quintile, then display post-subclassification balance	45
Check Balance and Propensity Score Overlap in Each Quintile	45
Creating a Standardized Difference Calculation Function	47
Creating the Five Subsamples, by PS Quintile	47
Standardized Differences in Each Quintile, and Overall	48
Plotting the Standardized Differences	48
Checking Rubin’s Rules Post-Subclassification	50
Task 7. After subclassifying, what is the estimated average treatment effect?	53
... on Outcome 1 [a continuous outcome]	53
... on Outcome 2 [a binary outcome]	54
... on Outcome 3 [a time to event]	55
Results So Far (After Matching and Subclassification)	56
Task 8. Execute weighting by the inverse PS, then assess covariate balance	57
ATT approach: Weight treated subjects as 1; control subjects as $ps/(1-ps)$	57
ATE Approach: Weight treated subjects by $1/ps$; Control subjects by $1/(1-PS)$	58
Assessing Balance after Weighting	59
Rubin’s Rules after ATT weighting	64
Rubin’s Rules after ATE weighting	65
Using TWANG for Alternative PS Estimation and ATT Weighting	65
Estimate the Propensity Score using Generalized Boosted Regression, and then perform ATT Weighting	65
Semi-Automated Love plot of Standardized Differences	69
Semi-Automated Love plot of Variance Ratios	69
Task 9. After weighting, what is the estimated average causal effect of treatment?	70
... on Outcome 1 [a continuous outcome]	70
... on Outcome 2 [a binary outcome]	71
... on Outcome 3 [a time to event]	71
Results So Far (After Matching, Subclassification and Weighting)	73
Task 10. After direct adjustment for the linear PS, what is the estimated average causal treatment effect?	74
... on Outcome 1 [a continuous outcome]	74
... on Outcome 2 [a binary outcome]	74
... on Outcome 3 [a time-to-event outcome]	74
Results So Far (After Matching, Subclassification, Weighting, Adjustment)	76
Task 11. “Double Robust” Approach - Weighting + Adjustment, what is the estimated average causal effect of treatment?	77
... on Outcome 1 [a continuous outcome]	77
... on Outcome 2 [a binary outcome]	78
... on Outcome 3 [a time to event]	79
Task 12. Results	82
Treatment Effect Estimates	82
Quality of Balance: Standardized Differences and Variance Ratios	83
Quality of Balance: Rubin’s Rules	83
What is a Sensitivity Analysis for Matched Samples?	84
Goal of a Formal Sensitivity Analysis for Matched Samples	84
The Sensitivity Parameter, Γ	84
Interpreting the Sensitivity Parameter, Γ	85

Task 13. Sensitivity Analysis for Matched Samples, Outcome 1, using rbounds	85
Rosenbaum Bounds for the Wilcoxon Signed Rank test (Quantitative outcome)	85
Specifying The Threshold Γ value	86
Interpreting Γ appropriately	87
Alternative Descriptions of Γ	87
An Alternate Approach - the Hodges-Lehman estimate	87
What about other types of outcomes?	88
What about when we match 1:2 or 1:3 instead of 1:1?	88
Wrapup	88

Setup

```
library(skimr)
library(tableone)
library(broom)
library(Epi)
library(survival)
library(Matching)
library(cobalt)
library(lme4)
library(twang)
library(survey)
library(rbounds)
library(tidyverse)

skim_with(numeric = list(hist = NULL),
          integer = list(hist = NULL))

decim <- function(x, k) format(round(x, k), nsmall=k)
```

The Data Set

The Data Set is 100% fictional, and is available as `toy2019.csv` on the course website.

- It contains data on 400 subjects (140 treated and 260 controls) on treatment status, six covariates, and three outcomes, with no missing observations anywhere.
- We assume that a logical argument suggests that the square of `covA`, as well as the interactions of `covB` with `covC` and with `covD` should be related to treatment assignment, and thus should be included in our propensity model.
- Our objective is to estimate the average causal effect of treatment (as compared to control) on each of the three outcomes, without propensity adjustment, and then with propensity matching, subclassification, weighting and regression adjustment using the propensity score.

```
toy <- read.csv("toy2019.csv") %>% tbl_df()

toy
```

```
# A tibble: 400 x 11
  subject treated covA covB covC covD covE covF out1.cost out2.event
  <fct>      <int> <dbl> <int> <dbl> <dbl> <int> <fct>      <int> <fct>
1 T_001         0  4.13     0  10     8.2    14 2-Mi~        34 No
2 T_002         1  4.58     1  8.69  10.1    13 1-Low        63 No
```

```

3 T_003      0  1.28      0 11.8      5.6      14 1-Low      61 No
4 T_004      0  3.11      0 10.9     10.9      10 1-Low      34 No
5 T_005      1  3.31      0 10.5      9         14 3-Hi~      38 No
6 T_006      0  4.08      0 13.9     10         5 3-Hi~      51 No
7 T_007      0  3.86      1 13        6.1      10 3-Hi~      53 Yes
8 T_008      0  2.58      0 12.6     5.2       4 2-Mi~      53 Yes
9 T_009      0  3.46      0 10.1     8.8      10 1-Low      61 Yes
10 T_010     0  3.11      1 13.3     4.8      10 1-Low      28 No
# ... with 390 more rows, and 1 more variable: out3.time <int>

```

The Codebook for the toy data

```

toy.codebook <- data_frame(
  Variable = dput(names(toy)),
  Type = c("Subject ID", "2-level categorical (0/1)", "Quantitative (2 decimal places)",
           "2-level categorical (0/1)", "Quantitative (1 decimal place)",
           "Quantitative (1 decimal place)", "Integer",
           "3-level ordinal factor", "Quantitative outcome",
           "Binary outcome (did event occur?)", "Time to event outcome"),
  Notes = c("labels are T_001 to T_400", "0 = control, 1 = treated",
            "reasonable values range from 0 to 6", "0 = no, 1 = yes",
            "plausible range 3-20", "plausible range 3-20", "plausible range 3-20",
            "1 = Low, 2 = Middle, 3 = High",
            "typical values 10-100", "Yes/No (note: event is bad)",
            "Time before event is observed or subject exits study (censored), range is 76-154 weeks")
)

c("subject", "treated", "covA", "covB", "covC", "covD", "covE",
  "covF", "out1.cost", "out2.event", "out3.time")

toy.codebook

```

```

# A tibble: 11 x 3
  Variable   Type      Notes
  <chr>      <chr>      <chr>
1 subject   Subject ID labels are T_001 to T_400
2 treated   2-level categorical 0 = control, 1 = treated
3 covA      Quantitative (2 de~ reasonable values range from 0 to 6
4 covB      2-level categorical 0 = no, 1 = yes
5 covC      Quantitative (1 de~ plausible range 3-20
6 covD      Quantitative (1 de~ plausible range 3-20
7 covE      Integer      plausible range 3-20
8 covF      3-level ordinal fa~ 1 = Low, 2 = Middle, 3 = High
9 out1.cost Quantitative outco~ typical values 10-100
10 out2.event Binary outcome (di~ Yes/No (note: event is bad)
11 out3.time Time to event outc~ Time before event is observed or subjec~

```

With regard to the out3.time variable, subjects with out2.event = No were censored, so that out2.event = Yes indicates an observed event.

“Skimmed” Summaries, within treatment groups

```
toy %>% group_by(treated) %>% skim(-subject)
```

```
Skim summary statistics
n obs: 400
n variables: 11
group variables: treated
```

```
-- Variable type:factor -----
```

treated	variable	missing	complete	n	n_unique
0	covF	0	260	260	3
0	out2.event	0	260	260	2
1	covF	0	140	140	3
1	out2.event	0	140	140	2

```
top_counts ordered
```

```
1-L: 118, 2-M: 98, 3-H: 44, NA: 0 FALSE
    No: 154, Yes: 106, NA: 0 FALSE
2-M: 54, 3-H: 48, 1-L: 38, NA: 0 FALSE
    Yes: 82, No: 58, NA: 0 FALSE
```

```
-- Variable type:integer -----
```

treated	variable	missing	complete	n	mean	sd	p0	p25	p50	p75
0	covB	0	260	260	0.3	0.46	0	0	0	1
0	covE	0	260	260	11.3	3.42	4	9	11	13.25
0	out1.cost	0	260	260	47.01	12.39	20	38	47	54
0	out3.time	0	260	260	109.85	12.61	79	101	110	118.25
1	covB	0	140	140	0.51	0.5	0	0	1	1
1	covE	0	140	140	9.77	2.84	4	8	9	12
1	out1.cost	0	140	140	56.64	16.56	20	45	56.5	72.25
1	out3.time	0	140	140	102.71	11.99	76	95	101	110

```
p100
```

```
1
```

```
19
```

```
84
```

```
154
```

```
1
```

```
16
```

```
84
```

```
136
```

```
-- Variable type:numeric -----
```

treated	variable	missing	complete	n	mean	sd	p0	p25	p50	p75
0	covA	0	260	260	3	1.09	0.2	2.51	3.08	3.84
0	covC	0	260	260	10.6	2.05	5.56	9.24	10.6	12.33
0	covD	0	260	260	8.65	2.21	2.8	7.2	9.05	10.3
1	covA	0	140	140	3.16	1.14	0.65	2.45	3.29	4.16
1	covC	0	140	140	9.62	1.87	5.96	8.17	9.58	10.8
1	covD	0	140	140	9.16	2.08	3.2	7.65	9.35	10.8

```
p100
```

```
5.35
```

```
14.44
```

```
12.8
```

```
5.05
```

```
13.94
```

```
14.5
```

Table 1

```
factorlist <- c("covB", "covF", "out2.event")

CreateTableOne(data = toy,
  vars = dput(names(select(toy, -subject, -treated))),
  strata = "treated", factorVars = factorlist)

c("covA", "covB", "covC", "covD", "covE", "covF", "out1.cost",
  "out2.event", "out3.time")
```

	Stratified by treated		p	test
	0	1		
n	260	140		
covA (mean (sd))	3.00 (1.09)	3.16 (1.14)	0.170	
covB = 1 (%)	77 (29.6)	72 (51.4)	<0.001	
covC (mean (sd))	10.60 (2.05)	9.62 (1.87)	<0.001	
covD (mean (sd))	8.65 (2.21)	9.16 (2.08)	0.025	
covE (mean (sd))	11.30 (3.42)	9.77 (2.84)	<0.001	
covF (%)			<0.001	
1-Low	118 (45.4)	38 (27.1)		
2-Middle	98 (37.7)	54 (38.6)		
3-High	44 (16.9)	48 (34.3)		
out1.cost (mean (sd))	47.01 (12.39)	56.64 (16.56)	<0.001	
out2.event = Yes (%)	106 (40.8)	82 (58.6)	0.001	
out3.time (mean (sd))	109.85 (12.61)	102.71 (11.99)	<0.001	

Data Management and Cleanup

Range Checks for Quantitative (continuous) Variables

Checking and cleaning the quantitative variables is pretty straightforward - the main thing I'll do at this stage is check the ranges of values shown to ensure that they match up with what I'm expecting. Here, all of the quantitative variables have values that fall within the “permissible” range described by my codebook, so we'll assume that for the moment, we're OK on `subject` (just a meaningless code, really), `covA`, `covC`, `covD`, `covE`, `out1.cost` and `out3.time`, and we see no missingness.

Restating Categorical Information in Helpful Ways

The cleanup of the toy data focuses, as it usually does, on variables that contain **categories** of information, rather than simple counts or measures, represented in quantitative variables.

Re-expressing Binary Variables as Numbers and Factors

We have three binary variables (`treated`, `covB` and `out2.event`). A major issue in developing these variables is to ensure that the direction of resulting odds ratios and risk differences are consistent and that cross-tabulations are in standard epidemiological format.

It will be useful to define binary variables in two ways:

- as a numeric indicator variable taking on the values 0 (meaning “not having the characteristic being studied”) or 1 (meaning “having the characteristic being studied”)

- as a text factor - with the levels of our key exposure and outcomes arranged so that “having the characteristic” precedes “not having the characteristic” in R when you create a table, but the covariates should still be No/Yes.

So what do we currently have? From the output below, it looks like `treated` and `covB` are numeric, 0/1 variables, while `out2.event` is a factor with levels “No” and then “Yes”

```
toy %>% select(treated, covB, out2.event) %>% summary()
```

treated	covB	out2.event
Min. :0.00	Min. :0.0000	No :212
1st Qu.:0.00	1st Qu.:0.0000	Yes:188
Median :0.00	Median :0.0000	
Mean :0.35	Mean :0.3725	
3rd Qu.:1.00	3rd Qu.:1.0000	
Max. :1.00	Max. :1.0000	

So, we’ll create factors for `treated` and `covB`:

```
toy$treated_f <- factor(toy$treated, levels = c(1,0),
                        labels = c("Treated", "Control"))
toy$covB_f <- factor(toy$covB, levels = c(0,1),
                    labels = c("No B", "Has B"))
```

For `out2.event`, on the other hand, we don’t have either quite the way we might want it. As you see in the summary output, we have two codes for `out2.event` - either No or Yes, in that order. But we want Yes to precede No (and I’d like a more meaningful name). So I redefine the factor variable, as follows.

```
toy$out2_f <- factor(toy$out2.event, levels = c("Yes","No"),
                    labels = c("Event","No Event"))
```

To obtain a numerical (0 or 1) version of `out2.event` we can use R’s `as.numeric` function - the problem is that this produces values of 1 (for No) and 2 (for Yes), rather than 0 and 1. So, I simply subtract 1 from the result, and we get what we need.

```
toy$out2 <- as.numeric(toy$out2.event) - 1
```

Testing Your Code - Sanity Checks

Before I move on, I’ll do a series of sanity checks to make sure that our new variables are defined as we want them, by producing a series of small tables comparing the new variables to those originally included in the data set.

```
toy %>% count(treated, treated_f)
```

```
# A tibble: 2 x 3
  treated treated_f    n
  <int> <fct>    <int>
1      0 Control    260
2      1 Treated    140
```

```
toy %>% count(covB, covB_f)
```

```
# A tibble: 2 x 3
  covB covB_f    n
  <int> <fct>    <int>
1      0 No B    251
2      1 Has B   149
```

```
toy %>% count(out2.event, out2_f, out2)
```

```
# A tibble: 2 x 4
  out2.event out2_f    out2     n
  <fct>      <fct>    <dbl> <int>
1 No        No Event      0    212
2 Yes       Event      1    188
```

Everything looks OK:

- `treated_f` correctly captures the information in `treated`, with the label Treated above the label Control in the rows of the table, facilitating standard epidemiological format.
- `covB_f` also correctly captures the `covB` information, placing “Has B” last.
- `out2_f` correctly captures and re-orders the labels from the original `out2.event`
- `out2` shows the data correctly (as compared to the original `out2.event`) with 0-1 coding.

Dealing with Variables including More than Two Categories

When we have a multi-categorical (more than two categories) variable, like `covF`, we will want to have

- both a text version of the variable with sensibly ordered levels, as a factor in R, as well as
- a series of numeric indicator variables (taking the values 0 or 1) for the individual levels.

```
toy %>% count(covF)
```

```
# A tibble: 3 x 2
  covF      n
  <fct>  <int>
1 1-Low   156
2 2-Middle 152
3 3-High   92
```

From the `summary` output, we can see that we’re all set for the text version of `covF`, as what we have currently is a factor with three levels, labeled 1-Low, 2-Middle and 3-High. This list of variables should work out well for us, as it preserves the ordering in a table and permits us to see the names, too. If we’d used just Low, Middle and High, then when R sorted a table into alphabetical order, we’d have High, then Low, then Middle - not ideal.

Preparing Indicator Variables for `covF`

So, all we need to do for `covF` is prepare indicator variables. We can either do this for all levels, or select one as the baseline, and do the rest. Here, I’ll show them all.

```
toy <- toy %>%
  mutate(covF.Low = as.numeric(covF == "1-Low"),
         covF.Middle = as.numeric(covF == "2-Middle"),
         covF.High = as.numeric(covF == "3-High"))
```

And now, some more sanity checks for the `covF` information:

```
toy %>% count(covF, covF.High, covF.Middle, covF.Low)
```

```
# A tibble: 3 x 5
  covF    covF.High covF.Middle covF.Low     n
  <fct>    <dbl>      <dbl>    <dbl> <int>
1 1-Low      0          0          1    156
```


2	2-Middle	0	1	0	152
3	3-High	1	0	0	92

Creating the Transformation and Product Terms

Remember that we have reason to believe that the square of `covA` as well as the interaction of `covB` with `covC` and also `covB` with `covD` will have an impact on treatment assignment. It will be useful to have these transformations in our data set for modeling and summarizing. I will use `covB` in its numeric (0,1) form (rather than as a factor - `covB.f`) when creating product terms, as shown below.

```
toy <- toy %>%
  mutate(Asqr = covA^2,
         BC = covB*covC,
         BD = covB*covD)
```

Data Set After Cleaning

Skim, within Treatment Groups

```
toy %>% select(treated_f, covA, covB, covC, covD, covE,
              covF, Asqr, BC, BD, out1.cost, out2, out3.time) %>%
  group_by(treated_f) %>%
  skim()
```

Skim summary statistics

n obs: 400

n variables: 13

group variables: treated_f

```
-- Variable type:factor -----
treated_f variable missing complete  n n_unique
Treated   covF           0       140 140         3
Control   covF           0       260 260         3
              top_counts ordered
2-M: 54, 3-H: 48, 1-L: 38, NA: 0  FALSE
1-L: 118, 2-M: 98, 3-H: 44, NA: 0  FALSE
```

```
-- Variable type:integer -----
treated_f variable missing complete  n  mean   sd p0 p25  p50  p75
Treated   covB           0       140 140   0.51  0.5  0  0  1  1
Treated   covE           0       140 140   9.77  2.84  4  8  9  12
Treated   out1.cost       0       140 140  56.64 16.56 20 45 56.5 72.25
Treated   out3.time       0       140 140 102.71 11.99 76 95 101 110
Control   covB           0       260 260   0.3   0.46  0  0  0  1
Control   covE           0       260 260  11.3   3.42  4  9  11 13.25
Control   out1.cost       0       260 260  47.01 12.39 20 38 47  54
Control   out3.time       0       260 260 109.85 12.61 79 101 110 118.25
p100
1
16
84
136
```

1
19
84
154

```
-- Variable type:numeric -----
treated_f variable missing complete  n  mean  sd  p0  p25  p50  p75
Treated    Asqr      0    140 140 11.3  6.74 0.42 6    10.82 17.26
Treated     BC      0    140 140  4.95  5.02 0     0     6.43  9.69
Treated     BD      0    140 140  4.52  4.66 0     0     4.25  9.2
Treated    covA      0    140 140  3.16  1.14 0.65 2.45  3.29  4.16
Treated    covC      0    140 140  9.62  1.87 5.96 8.17  9.58 10.8
Treated    covD      0    140 140  9.16  2.08 3.2   7.65  9.35 10.8
Treated    out2      0    140 140  0.59  0.49 0     0     1     1
Control    Asqr      0    260 260 10.22  6.01 0.04 6.3   9.49 14.78
Control     BC      0    260 260  2.99  4.78 0     0     0     7.38
Control     BD      0    260 260  2.44  3.93 0     0     0     6.1
Control    covA      0    260 260  3     1.09 0.2   2.51  3.08  3.84
Control    covC      0    260 260 10.6   2.05 5.56 9.24 10.6 12.33
Control    covD      0    260 260  8.65  2.21 2.8   7.2   9.05 10.3
Control    out2      0    260 260  0.41  0.49 0     0     0     1
p100
25.5
13.7
12.2
5.05
13.94
14.5
1
28.62
14.24
12.5
5.35
14.44
12.8
1
```

Table 1

Note that the factors I created for the `out2` outcome are not well ordered for a Table 1, but are well ordered for other tables we'll fit later. So, in this case, I'll use the numeric version of the `out2` outcome, but the new factor representations of `covB` and `treated`.

```
varlist = c("covA", "covB_f", "covC", "covD", "covE", "covF",
            "Asqr", "BC", "BD", "out1.cost", "out2", "out3.time")
factorlist = c("covB_f", "covF", "out2")
CreateTableOne(vars = varlist, strata = "treated_f",
               data = toy, factorVars = factorlist)
```

	Stratified by treated_f			
	Treated	Control	p	test
n	140	260		
covA (mean (sd))	3.16 (1.14)	3.00 (1.09)	0.170	
covB_f = Has B (%)	72 (51.4)	77 (29.6)	<0.001	

covC (mean (sd))	9.62 (1.87)	10.60 (2.05)	<0.001
covD (mean (sd))	9.16 (2.08)	8.65 (2.21)	0.025
covE (mean (sd))	9.77 (2.84)	11.30 (3.42)	<0.001
covF (%)			<0.001
1-Low	38 (27.1)	118 (45.4)	
2-Middle	54 (38.6)	98 (37.7)	
3-High	48 (34.3)	44 (16.9)	
Asqr (mean (sd))	11.30 (6.74)	10.22 (6.01)	0.101
BC (mean (sd))	4.95 (5.02)	2.99 (4.78)	<0.001
BD (mean (sd))	4.52 (4.66)	2.44 (3.93)	<0.001
out1.cost (mean (sd))	56.64 (16.56)	47.01 (12.39)	<0.001
out2 = 1 (%)	82 (58.6)	106 (40.8)	0.001
out3.time (mean (sd))	102.71 (11.99)	109.85 (12.61)	<0.001

The 13 Tasks We'll Tackle in this Example

1. Ignoring the covariate information, what is the unadjusted point estimate (and 95% confidence interval) for the effect of the treatment on each of the three outcomes (`out1.cost`, `out2.event`, and `out3.time`)?
2. Assume that theory suggests that the square of `covA`, as well as the interactions of `covB` with `covC` and `covB` with `covD` should be related to treatment assignment. Fit a propensity score model to the data, using the six covariates (A-F) and the three transformations (A^2 , and the B-C and B-D interactions.) Plot the resulting propensity scores, by treatment group, in an attractive and useful way.
3. Use Rubin's Rules to assess the overlap of the propensity scores and the individual covariates prior to the use of any propensity score adjustments.
4. Use 1:1 greedy matching to match all 140 treated subjects to control subjects without replacement on the basis of the linear propensity for treatment. Evaluate the degree of covariate imbalance before and after propensity matching for each of the six covariates, and present the pre- and post-match standardized differences and variance ratios for the covariates, as well as the square term and interactions, as well as both the raw and linear propensity score in appropriate plots. Now, build a new data frame containing the propensity-matched sample, and use it to first check Rubin's Rules after matching.
5. Now, use the matched sample data set to evaluate the treatment's average causal effect on each of the three outcomes. In each case, specify a point estimate (and associated 95% confidence interval) for the effect of being treated (as compared to being a control subject) on the outcome. Compare your results to the automatic versions reported by the Matching package when you include the outcome in the matching process.
6. Now, instead of matching, instead subclassify the subjects into quintiles by the raw propensity score. Display the balance in terms of standardized differences by quintile for the covariates, their transformations, and the propensity score in an appropriate table or plot(s). Are you satisfied?
7. Regardless of your answer to the previous question, use the propensity score quintile subclassification approach to find a point estimate (and 95% confidence interval) for the effect of the treatment on each outcome.
8. Now using a reasonable propensity score weighting strategy, assess the balance of each covariate, the transformations and the linear propensity score prior to and after propensity weighting. Is the balance after weighting satisfactory?
9. Using propensity score weighting to evaluate the treatment's effect, developing a point estimate and 95% CI for the average causal effect of treatment on each outcome.
10. Finally, use direct adjustment for the linear propensity score on the entire sample to evaluate the treatment's effect, developing a point estimate and 95% CI for each outcome.
11. Now, try a double robust approach. Weight, then adjust for linear propensity score.
12. Compare your conclusions about the average causal effect obtained in the following six ways to each other. What happens and why? Which of these methods seems most appropriate given the available information?

- without propensity adjustment,
- after propensity matching,
- after propensity score subclassification,
- after propensity score weighting,
- after adjusting for the propensity score directly, and
- after weighting then adjusting for the PS, to each other.

13. Perform a sensitivity analysis for your matched samples analysis and the first outcome (`out1.cost`) if it turns out to show a statistically significant treatment effect.

Task 1. Ignoring covariates, estimate the effect of treatment vs. control on...

Outcome 1 (a continuous outcome)

Our first outcome describes a quantitative measure, cost, and we're asking what the effect of `treatment` as compared to `control` is on that outcome. Starting with brief numerical summaries:

```
toy %>%
  group_by(treated_f) %>%
  skim(out1.cost)
```

Skim summary statistics

n obs: 400

n variables: 21

group variables: treated_f

```
-- Variable type:integer -----
treated_f variable missing complete  n  mean    sd p0 p25  p50  p75
Treated out1.cost          0      140 140 56.64 16.56 20  45 56.5 72.25
Control out1.cost          0      260 260 47.01 12.39 20  38 47  54
p100
84
84
```

It looks like the Treated group has higher costs than the Control group. To model this, we could use a linear regression model to obtain a point estimate and 95% confidence interval. Here, I prefer to use the numeric version of the `treated` variable, with 0 = “control” and 1 = “treated”.

```
unadj.out1 <- lm(out1.cost ~ treated, data=toy)
summary(unadj.out1); confint(unadj.out1, level = 0.95) ## provides treated effect and CI estimates
```

Call:

```
lm(formula = out1.cost ~ treated, data = toy)
```

Residuals:

Min	1Q	Median	3Q	Max
-36.643	-11.008	-0.008	9.084	36.992

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	47.0077	0.8673	54.202	< 2e-16 ***
treated	9.6352	1.4659	6.573	1.55e-10 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 13.98 on 398 degrees of freedom
Multiple R-squared:  0.09791,    Adjusted R-squared:  0.09565
F-statistic: 43.2 on 1 and 398 DF,  p-value: 1.553e-10
```

```
          2.5 %    97.5 %
(Intercept) 45.302702 48.71268
treated      6.753205 12.51713
```

We can store these results in a data frame, with the `tidy` function from the `broom` package.

```
tidy(unadj.out1, conf.int = TRUE, conf.level = 0.95)
```

```
# A tibble: 2 x 7
  term      estimate std.error statistic  p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 (Intercept)  47.0       0.867     54.2 7.71e-186  45.3     48.7
2 treated       9.64      1.47      6.57 1.55e-10   6.75    12.5
```

```
res_unadj_1 <- tidy(unadj.out1, conf.int = TRUE, conf.level = 0.95) %>%
  filter(term == "treated")
```

```
res_unadj_1
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic  p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated       9.64      1.47      6.57 1.55e-10   6.75    12.5
```

Our unadjusted treatment effect estimate is a difference of 9.64 in cost, with 95% confidence interval (6.75, 12.52).

Outcome 2 (a binary outcome)

Using a 2x2 table in standard epidemiological format

Thanks to our preliminary cleanup, it's relatively easy to obtain a table in standard epidemiological format comparing treated to control subjects in terms of `out2`:

```
table(toy$treated_f, toy$out2_f)
```

	Event	No Event
Treated	82	58
Control	106	154

Note that the exposure is in the rows, with “Having the Exposure” or “Treated” at the top, and the outcome is in the columns, with “Yes” or “Outcome Occurred” or “Event Occurred” on the left, so that the top left cell count describes people that had both the exposure and the outcome. That's *standard epidemiological format*, just what we need for the `twoby2` function in the `Epi` package.

```
temp <- twoby2(table(toy$treated_f, toy$out2_f))
```

2 by 2 table analysis:

```
-----
Outcome   : Event
```

Comparing : Treated vs. Control

	Event	No Event	P(Event)	95% conf. interval	
Treated	82	58	0.5857	0.5025	0.6643
Control	106	154	0.4077	0.3496	0.4685

	95% conf. interval		
Relative Risk:	1.4367	1.1737	1.7586
Sample Odds Ratio:	2.0540	1.3530	3.1181
Conditional MLE Odds Ratio:	2.0502	1.3248	3.1884
Probability difference:	0.1780	0.0754	0.2754

Exact P-value: 8e-04

Asymptotic P-value: 7e-04

Eventually, we will be interested in at least two measures - the odds ratio and the risk (probability) difference estimates, and their respective confidence intervals.

The risk difference is shown as the Probability difference here. Let's save it to a data frame, and then we'll save the (sample) odds ratio information to another data frame.

```
res_unadj_2_riskdiff <- data.frame(out = "out2.event",  
  risk.diff = temp$measures[4,1],  
  conf.low = temp$measures[4,2],  
  conf.high = temp$measures[4,3])
```

```
res_unadj_2_oddsratio <- data.frame(out = "out2.event",  
  odds.ratio = temp$measures[2,1],  
  conf.low = temp$measures[2,2],  
  conf.high = temp$measures[2,3])
```

```
res_unadj_2_riskdiff
```

```
      out risk.diff  conf.low conf.high  
1 out2.event  0.178022 0.07536515 0.2753985
```

```
res_unadj_2_oddsratio
```

```
      out odds.ratio conf.low conf.high  
1 out2.event  2.054001 1.353022 3.118147
```

- For a *difference in risk*, our unadjusted treatment effect estimate is an difference of 17.8 percentage points as compared to control, with 95% CI of (7.5, 27.5) percentage points.
- For an *odds ratio*, our unadjusted treatment effect estimate is an odds ratio of 2.05 (95% CI = 1.35, 3.12) for the event occurring with treatment as compared to control.

Using a logistic regression model

For the odds ratio estimate, we can use a simple logistic regression model to estimate the unadjusted treatment effect, resulting in essentially the same answer. We'll use the numerical (0/1) format to represent binary information, as follows.

```
unadj.out2 <- glm(out2 ~ treated, data=toy, family=binomial())
```

```
summary(unadj.out2)
```

```
Call:
glm(formula = out2 ~ treated, family = binomial(), data = toy)
```

```
Deviance Residuals:
```

```
      Min       1Q   Median       3Q      Max
-1.328  -1.023  -1.023   1.340   1.340
```

```
Coefficients:
```

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -0.3735     0.1262  -2.960 0.003080 **
treated         0.7198     0.2130   3.379 0.000726 ***
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for binomial family taken to be 1)
```

```
Null deviance: 553.08  on 399  degrees of freedom
Residual deviance: 541.47  on 398  degrees of freedom
AIC: 545.47
```

```
Number of Fisher Scoring iterations: 4
```

```
exp(coef(unadj.out2)) # produces odds ratio estimate
```

```
(Intercept)      treated
  0.6883117    2.0540013
```

```
exp(confint(unadj.out2)) # produces 95% CI for odds ratio
```

```
              2.5 %    97.5 %
(Intercept) 0.5362913 0.8800944
treated      1.3561085 3.1283210
```

And, again, we can use the `tidy` function in the `broom` package to build a tibble of the key parts of the output. Note that by including the `exponentiate = TRUE` command, our results in the `treated` row describe the odds ratio, rather than the log odds.

```
tidy(unadj.out2, conf.int = TRUE, exponentiate = TRUE)
```

```
# A tibble: 2 x 7
```

```
  term          estimate std.error statistic  p.value conf.low conf.high
<chr>         <dbl>     <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 (Intercept)   0.688     0.126    -2.96 0.00308    0.536    0.880
2 treated       2.05      0.213     3.38 0.000726    1.36     3.13
```

```
res_unadj_2_or <- tidy(unadj.out2, conf.int = TRUE,
  conf.level = 0.95, exponentiate = TRUE) %>%
  filter(term == "treated")
```

```
res_unadj_2_or
```

```
# A tibble: 1 x 7
```

```
  term          estimate std.error statistic  p.value conf.low conf.high
<chr>         <dbl>     <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 treated       2.05      0.213     3.38 0.000726    1.36     3.13
```

- Our odds ratio estimate is 2.05, with 95% confidence interval ranging from 1.36 to 3.13.

- For practical purposes, the odds ratio and 95% confidence interval obtained here matches the methodology for the `twoby2` function. The approach implemented in the `twoby2` function produces slightly less conservative (i.e. narrower) confidence intervals for the effect estimate than does the approach used in the logistic regression model.

Outcome 3 (a time-to-event outcome with right censoring)

Our `out3.time` variable is a variable indicating the time before the event described in `out2` occurred. This happened to 188 of the 400 subjects in the data set. For the other 212 subjects who left the study before their event occurred, we have the time before censoring. We can see the results of this censoring in the survival object describing each treatment group.

Here, for instance, is the survival object for the *treated* subjects - the first subject listed here is censored - had the event at some point after 106 weeks (106+) but we don't know precisely when after 106 weeks.

```
Surv(toy$out3.time, toy$out2.event == "Yes")[toy$treated == 1]
```

```
[1] 106+ 96+ 96 99+ 99 108+ 124 116+ 101+ 110 80+ 94 99 126
[15] 93 93+ 104 125 102 87 99 102+ 101+ 101 83 94 107 130+
[29] 112+ 111 95 96 80+ 89 110 116+ 108+ 118 95 125+ 104 103
[43] 112+ 115+ 90 110+ 105+ 113+ 136+ 105 96+ 126+ 108+ 96 116+ 99
[57] 96 108+ 109 114 112 108+ 115 112+ 100 115+ 114+ 109 127+ 100
[71] 85 110 115 117 88 91 78+ 104+ 96+ 100+ 108+ 107+ 116 91
[85] 88 127+ 99 96+ 87 120+ 108 99 87 101 106+ 97 128 100
[99] 94 94 89 102 96 76 99+ 93 93 110 96+ 95 97 104
[113] 94 114+ 97+ 95 103+ 100+ 100 91 110+ 119 112+ 98 102+ 103
[127] 118+ 89 98+ 79 101+ 85 109+ 87 92 79+ 108+ 102 85 119+
```

- To see the controls, we could use `Surv(toy$out3.time, toy$out2.event=="Yes")[toy$treated==0]`

To deal with the right censoring, we'll use the `survival` package to fit a simple unadjusted Cox proportional hazards model to assess the relative hazard of having the event at a particular time point among treated subjects as compared to controls.

```
unadj.out3 <- coxph(Surv(out3.time, out2.event=="Yes") ~ treated, data=toy)
summary(unadj.out3) ## exp(coef) section indicates relative risk estimate and 95% CI
```

Call:

```
coxph(formula = Surv(out3.time, out2.event == "Yes") ~ treated,
      data = toy)
```

```
n= 400, number of events= 188
```

```
      coef exp(coef) se(coef)      z Pr(>|z|)
treated 0.7737    2.1677  0.1489 5.196 2.04e-07 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
      exp(coef) exp(-coef) lower .95 upper .95
treated    2.168    0.4613    1.619    2.902
```

```
Concordance= 0.6 (se = 0.019 )
```

```
Rsquare= 0.062 (max possible= 0.993 )
```

```
Likelihood ratio test= 25.63 on 1 df, p=4e-07
```

```
Wald test            = 27 on 1 df, p=2e-07
```

```
Score (logrank) test = 28.3 on 1 df, p=1e-07
```


The relative hazard rate is shown in the `exp(coef)` section of the output. Our unadjusted treatment model suggests that the hazard of the outcome is smaller (but not significantly smaller) in the treated group than in the control group. Our estimate is that this relative hazard rate for occurrence of the event associated with treatment as compared to control is 0.86 with a 95% confidence interval of (0.57, 1.29).

Yes, you can tidy this model, as well, using the `broom` package.

```
res_unadj_3 <- tidy(unadj.out3, exponentiate = TRUE) %>%  
  filter(term == "treated")  
res_unadj_3
```

```
# A tibble: 1 x 7  
  term      estimate std.error statistic    p.value conf.low conf.high  
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>  
1 treated      2.17      0.149      5.20 0.000000204      1.62      2.90
```

And so, our estimate can be saved, as we've done previously.


- The relative hazard rate estimate is 2.17, with 95% confidence interval ranging from 1.62 to 2.90.

It's wise, whenever fitting a Cox proportional hazards model, to assess the proportional hazards assumption. One way to do this is to run a simple test in R - from which we can obtain a plot, if we like. The idea is for the plot to show no clear patterns over time, and look pretty much like a horizontal line, while we would like the test to be non-significant - if that's the case, our proportional hazards assumption is likely OK.

```
cox.zph(unadj.out3)
```

```
      rho chisq      p  
treated -0.0783  1.12 0.289
```

```
plot(cox.zph(unadj.out3), var="treated")
```



If the proportional hazards assumption is clearly violated (here it isn't), call a statistician.

Unadjusted Estimates of Treatment Effect on Outcomes

So, our unadjusted average treatment effect estimates (in each case comparing treated subjects to control subjects) are thus:

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Relative Hazard Rate)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)

Task 2. Fit the propensity score model, then plot the PS-treatment relationship

I'll use a logistic regression model

```
psmodel <- glm(treated ~ covA + covB + covC + covD + covE + covF +
               Asqr + BC + BD, family=binomial(), data=toy)
summary(psmodel)
```

```
Call:
glm(formula = treated ~ covA + covB + covC + covD + covE + covF +
     Asqr + BC + BD, family = binomial(), data = toy)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.7156	-0.8403	-0.5277	1.0302	1.9581

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.55195	1.54168	1.655	0.09786 .
covA	-0.31630	0.45711	-0.692	0.48896
covB	-1.64510	1.85012	-0.889	0.37390
covC	-0.26162	0.08627	-3.033	0.00243 **
covD	0.06869	0.07988	0.860	0.38986
covE	-0.15560	0.03943	-3.947	7.93e-05 ***
covF2-Middle	0.23060	0.27497	0.839	0.40167
covF3-High	0.90026	0.30555	2.946	0.00322 **
Asqr	0.07081	0.08095	0.875	0.38169
BC	0.22538	0.12432	1.813	0.06984 .
BD	0.04450	0.11894	0.374	0.70829

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 517.96 on 399 degrees of freedom
 Residual deviance: 444.25 on 389 degrees of freedom
 AIC: 466.25

Number of Fisher Scoring iterations: 4

Having fit the model, my first step will be to save the raw and linear propensity score values to the main toy example tibble.

```
toy$ps <- psmodel$fitted
toy$linps <- psmodel$linear.predictors
```

Comparing the Distribution of Propensity Score Across the Two Treatment Groups

Now, I can use these saved values to assess the propensity model.

```
toy %>% group_by(treated_f) %>% skim(ps, linps)
```

Skim summary statistics


n obs: 400
 n variables: 23
 group variables: treated_f

```
-- Variable type:numeric -----
treated_f variable missing complete n mean sd p0 p25 p50
Treated linps 0 140 140 -0.19 0.8 -1.76 -0.82 -0.15
```


Treated	ps	0	140	140	0.46	0.18	0.15	0.31	0.46
Control	linps	0	260	260	-1.08	1.01	-3.33	-1.8	-1.14
Control	ps	0	260	260	0.29	0.18	0.034	0.14	0.24
p75 p100									
0.43 1.61									
0.61 0.83									
-0.38 1.21									
0.4 0.77									

The simplest plot is probably a boxplot, but it's not very granular.

```
ggplot(toy, aes(x = treated_f, y = ps)) +  
  geom_boxplot()
```




```
ggplot(toy, aes(x = treated_f, y = ps, color = treated_f)) +  
  geom_boxplot() +  
  geom_jitter(width = 0.1) +  
  guides(color = FALSE)
```



I'd rather get a fancier plot to compare the distributions of the propensity score across the two treatment groups, perhaps using a smoothed density estimate, as shown below. Here, I'll show the distributions of the linear propensity score, the log odds of treatment.

```
ggplot(toy, aes(x = linps, fill = treated_f)) +  
  geom_density(alpha = 0.3)
```



We see a fair amount of overlap across the two treatment groups. I'll use Rubin's Rules in the next section to help assess the amount of overlap at this point, before any adjustments for the propensity score.

Task 3. Rubin's Rules to Check Overlap Before Propensity Adjustment

In his 2001 article¹ about using propensity scores to design studies, as applied to studies of the causal effects of the conduct of the tobacco industry on medical expenditures, Donald Rubin proposed three “rules” for assessing the overlap / balance of covariates appropriately before and after propensity adjustment. Before an outcome is evaluated using a regression analysis (perhaps supplemented by a propensity score adjustment through matching, weighting, subclassification or even direct adjustment), there are three checks that should be performed.

When we do a propensity score analysis, it will be helpful to perform these checks as soon as the propensity model has been estimated, even before any adjustments take place, to see how well the distributions of covariates overlap. After using the propensity score, we hope to see these checks meet the standards below. In what follows, I will describe each standard, and demonstrate its evaluation using the propensity score model we just fit, and looking at the original toy data set, without applying the propensity score in any way to do adjustments.

¹Rubin DB 2001 Using Propensity Scores to Help Design Observational Studies: Application to the Tobacco Litigation. *Health Services & Outcomes Research Methodology* 2: 169-188.

Rubin's Rule 1

Rubin's Rule 1 states that the absolute value of the standardized difference of the linear propensity score, comparing the treated group to the control group, should be close to 0, ideally below 10%, and in any case less than 50%. If so, we may move on to Rule 2.

To evaluate this rule in the toy example, we'll run the following code to place the right value into a variable called `rubin1.unadj` (for Rubin's Rule 1, unadjusted).

```
rubin1.unadj <- with(toy,
  abs(100*(mean(linps[treated==1]) - mean(linps[treated==0]))/sd(linps)))
rubin1.unadj
```

```
[1] 85.85784
```

What this does is calculate the (absolute value of the) standardized difference of the linear propensity score comparing treated subjects to control subjects.

- We want this value to be close to 0, and certainly less than 50 in order to push forward to outcomes analysis without further adjustment for the propensity score.
- Clearly, here, with a value above 50%, we can't justify simply running an unadjusted regression model, be it a linear, logistic or Cox model - we've got observed selection bias, and need to actually apply the propensity score somehow in order to account for this.
- So, we'll need to match, subclassify, weight or directly adjust for propensity here.

Since we've failed Rubin's 1st Rule, in some sense, we're done checking the rules, because we clearly need to further adjust for observed selection bias - there's no need to prove that further through checking Rubin's 2nd and 3rd rules. But we'll do it here to show what's involved.

Rubin's Rule 2

Rubin's Rule 2 states that the ratio of the variance of the linear propensity score in the treated group to the variance of the linear propensity score in the control group should be close to 1, ideally between 4/5 and 5/4, but certainly not very close to or exceeding 1/2 and 2. If so, we may move on to Rule 3.

To evaluate this rule in the toy example, we'll run the following code to place the right value into a variable called `rubin2.unadj` (for Rubin's Rule 2, unadjusted).

```
rubin2.unadj <- with(toy, var(linps[treated==1])/var(linps[treated==0]))
rubin2.unadj
```

```
[1] 0.6274233
```

This is the ratio of variances of the linear propensity score comparing treated subjects to control subjects. We want this value to be close to 1, and certainly between 0.5 and 2. In this case, we pass Rule 2, if just barely.

Rubin's Rule 3

For Rubin's Rule 3, we begin by calculating regression residuals for each covariate of interest (usually, each of those included in the propensity model) regressed on a single predictor - the linear propensity score. We then look to see if the ratio of the variance of the residuals of this model for the treatment group divided by the variance of the residuals of this model for the control group is close to 1. Again, ideally this will fall between 4/5 and 5/4 for each covariate, but certainly between 1/2 and 2. If so, then the use of regression models seems well justified.

To evaluate Rubin's 3rd Rule, we'll create a little function to help us do the calculations.

```
## General function rubin3 to help calculate Rubin's Rule 3
rubin3 <- function(data, covlist, linps) {
  covlist2 <- as.matrix(covlist)
  res <- NA
  for(i in 1:ncol(covlist2)) {
    cov <- as.numeric(covlist2[,i])
    num <- var(resid(lm(cov ~ data$linps)))[data$exposure == 1]
    den <- var(resid(lm(cov ~ data$linps)))[data$exposure == 0]
    res[i] <- decim(num/den, 3)
  }
  final <- data_frame(name = names(covlist), resid.var.ratio = as.numeric(res))
  return(final)
}
```

Now, then, applying the rule to our sample prior to propensity score adjustment, we get the following result. Note that I'm using the indicator variable forms for the covF information.

```
cov.sub <- toy %>% select(covA, covB, covC, covD, covE,
                        covF.Middle, covF.High, Asqr, BC, BD)

toy$exposure <- toy$treated

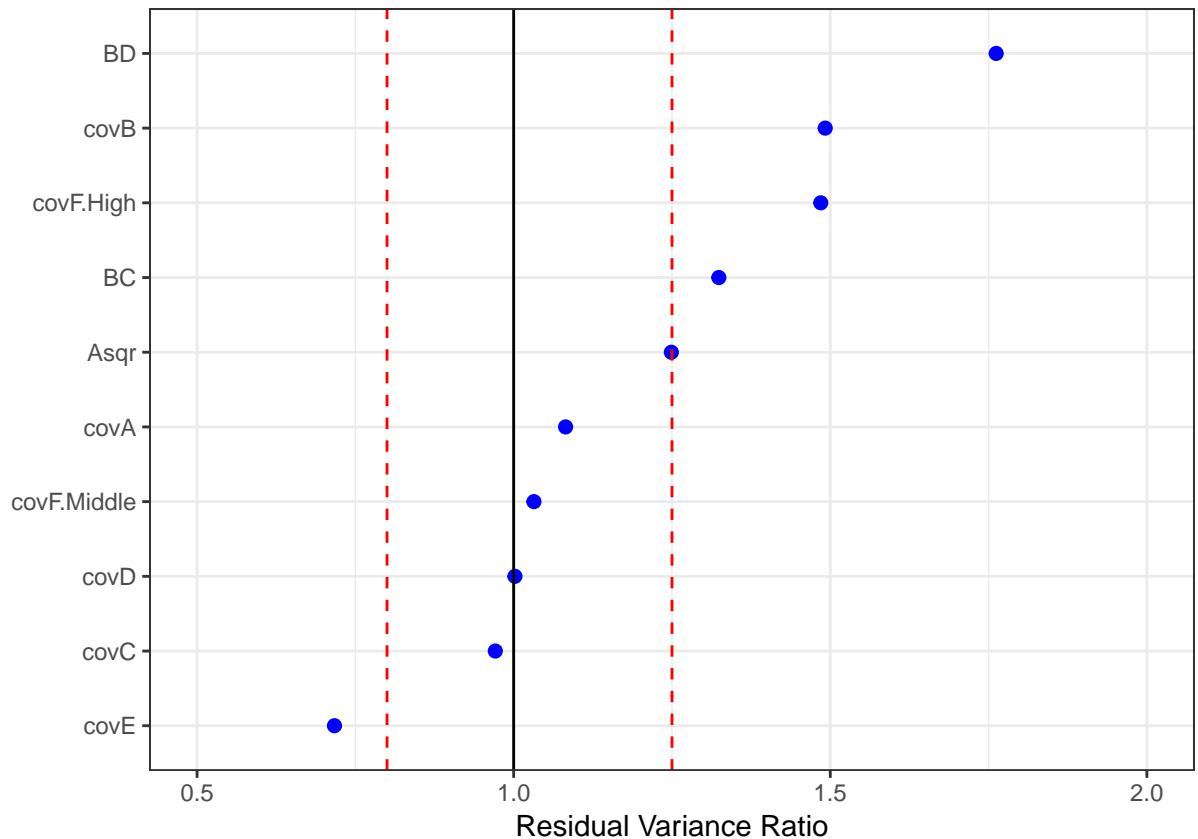
rubin3.unadj <- rubin3(data = toy, covlist = cov.sub, linps = linps)
rubin3.unadj
```

```
# A tibble: 10 x 2
  name      resid.var.ratio
  <chr>          <dbl>
1 covA          1.08
2 covB          1.49
3 covC          0.971
4 covD          1.00
5 covE          0.717
6 covF.Middle    1.03
7 covF.High     1.48
8 Asqr          1.25
9 BC            1.32
10 BD           1.76
```

Some of these covariates look to have residual variance ratios near 1, while others are further away, but all are within the (0.5, 2.0) range. So we'd pass Rule 3 here, although we'd clearly like to see some covariates (A and E, in particular) with ratios closer to 1.

A Cleveland Dot Chart of the Rubin's Rule 3 Results

```
ggplot(rubin3.unadj, aes(x = resid.var.ratio, y = reorder(name, resid.var.ratio))) +
  geom_point(col = "blue", size = 2) +
  theme_bw() +
  xlim(0.5, 2.0) +
  geom_vline(aes(xintercept = 1)) +
  geom_vline(aes(xintercept = 4/5), linetype = "dashed", col = "red") +
  geom_vline(aes(xintercept = 5/4), linetype = "dashed", col = "red") +
  labs(x = "Residual Variance Ratio", y = "")
```

We see values outside the $4/5$ and $5/4$ lines, but nothing falls outside $(0.5, 2)$.

Task 4. Use 1:1 greedy matching on the linear PS, then check post-match balance

As requested, we'll do 1:1 greedy matching on the linear propensity score without replacement and breaking ties randomly. To start, we won't include an outcome variable in our call to the `Match` function within the `Matching` package. We'll wind up with a match including 140 treated and 140 control subjects.

```
X <- toy$linps ## matching on the linear propensity score
Tr <- as.logical(toy$treated)
match1 <- Match(Tr=Tr, X=X, M = 1, replace=FALSE, ties=FALSE)
summary(match1)
```

```
Estimate... 0
SE..... 0
T-stat..... NaN
p.val..... NA
```

```
Original number of observations..... 400
Original number of treated obs..... 140
Matched number of observations..... 140
Matched number of observations (unweighted). 140
```

Balance Assessment (Semi-Automated)

Next, we'll assess the balance imposed by this greedy match on our covariates, and their transformations (A^2 and $B \cdot C$ and $B \cdot D$) as well as the raw and linear propensity scores. The default output from the `MatchBalance` function is extensive...

```
set.seed(5001)
mb1 <- MatchBalance(treated ~ covA + covB + covC + covD + covE + covF +
                    Asqr + BC + BD + ps + linps, data=toy,
                    match.out = match1, nboots=500)
```

***** (V1) covA *****

	Before Matching	After Matching
mean treatment.....	3.1646	3.1646
mean control.....	3.0046	3.0797
std mean diff.....	14.051	7.4541
mean raw eQQ diff.....	0.19193	0.15271
med raw eQQ diff.....	0.21	0.15
max raw eQQ diff.....	0.58	0.58
mean eCDF diff.....	0.047314	0.036354
med eCDF diff.....	0.035165	0.035714
max eCDF diff.....	0.11868	0.1
var ratio (Tr/Co).....	1.0837	1.0015
T-test p-value.....	0.1753	0.52808
KS Bootstrap p-value..	0.138	0.444
KS Naive p-value.....	0.154	0.48581
KS Statistic.....	0.11868	0.1

***** (V2) covB *****

	Before Matching	After Matching
mean treatment.....	0.51429	0.51429
mean control.....	0.29615	0.45
std mean diff.....	43.488	12.816
mean raw eQQ diff.....	0.22143	0.064286
med raw eQQ diff.....	0	0
max raw eQQ diff.....	1	1
mean eCDF diff.....	0.10907	0.032143
med eCDF diff.....	0.10907	0.032143
max eCDF diff.....	0.21813	0.064286
var ratio (Tr/Co).....	1.2023	1.0093
T-test p-value.....	2.6605e-05	0.20711

***** (V3) covC *****

	Before Matching	After Matching
mean treatment.....	9.6238	9.6238
mean control.....	10.596	9.7818

std mean diff.....	-51.896	-8.4375
mean raw eQQ diff.....	0.9755	0.20914
med raw eQQ diff.....	0.975	0.15
max raw eQQ diff.....	1.64	0.9
mean eCDF diff.....	0.12933	0.024845
med eCDF diff.....	0.13297	0.021429
max eCDF diff.....	0.24066	0.078571
var ratio (Tr/Co).....	0.83836	0.87377
T-test p-value.....	2.582e-06	0.46229
KS Bootstrap p-value..	< 2.22e-16	0.756
KS Naive p-value.....	5.2867e-05	0.7805
KS Statistic.....	0.24066	0.078571

***** (V4) covD *****

	Before Matching	After Matching
mean treatment.....	9.1593	9.1593
mean control.....	8.6469	9.2071
std mean diff.....	24.595	-2.2973
mean raw eQQ diff.....	0.54071	0.17929
med raw eQQ diff.....	0.5	0.1
max raw eQQ diff.....	1.8	1.7
mean eCDF diff.....	0.051117	0.01716
med eCDF diff.....	0.054945	0.014286
max eCDF diff.....	0.11648	0.05
var ratio (Tr/Co).....	0.8872	1.0941
T-test p-value.....	0.022381	0.84685
KS Bootstrap p-value..	0.128	0.98
KS Naive p-value.....	0.16916	0.9948
KS Statistic.....	0.11648	0.05

***** (V5) covE *****

	Before Matching	After Matching
mean treatment.....	9.7714	9.7714
mean control.....	11.3	10.05
std mean diff.....	-53.833	-9.8107
mean raw eQQ diff.....	1.5143	0.46429
med raw eQQ diff.....	2	0
max raw eQQ diff.....	4	2
mean eCDF diff.....	0.095673	0.035714
med eCDF diff.....	0.074725	0.0071429
max eCDF diff.....	0.22473	0.12857
var ratio (Tr/Co).....	0.68813	1.1133
T-test p-value.....	2.7506e-06	0.3651

KS Bootstrap p-value..	< 2.22e-16	0.088
KS Naive p-value.....	0.00020385	0.19748
KS Statistic.....	0.22473	0.12857

***** (V6) covF2-Middle *****

	Before Matching	After Matching
mean treatment.....	0.38571	0.38571
mean control.....	0.37692	0.44286
std mean diff.....	1.7996	-11.697
mean raw eQQ diff.....	0.0071429	0.057143
med raw eQQ diff.....	0	0
max raw eQQ diff.....	1	1
mean eCDF diff.....	0.0043956	0.028571
med eCDF diff.....	0.0043956	0.028571
max eCDF diff.....	0.0087912	0.057143
var ratio (Tr/Co).....	1.0122	0.9603
T-test p-value.....	0.86353	0.27615

***** (V7) covF3-High *****

	Before Matching	After Matching
mean treatment.....	0.34286	0.34286
mean control.....	0.16923	0.24286
std mean diff.....	36.448	20.992
mean raw eQQ diff.....	0.17143	0.1
med raw eQQ diff.....	0	0
max raw eQQ diff.....	1	1
mean eCDF diff.....	0.086813	0.05
med eCDF diff.....	0.086813	0.05
max eCDF diff.....	0.17363	0.1
var ratio (Tr/Co).....	1.6079	1.2253
T-test p-value.....	0.00023805	0.025801

***** (V8) Asqr *****

	Before Matching	After Matching
mean treatment.....	11.301	11.301
mean control.....	10.219	10.769
std mean diff.....	16.05	7.8879
mean raw eQQ diff.....	1.2406	0.87636
med raw eQQ diff.....	1.266	0.7181
max raw eQQ diff.....	3.2912	3.12
mean eCDF diff.....	0.047314	0.036354
med eCDF diff.....	0.035165	0.035714
max eCDF diff.....	0.11868	0.1

var ratio (Tr/Co).....	1.2571	1.1734
T-test p-value.....	0.11328	0.48574
KS Bootstrap p-value..	0.138	0.444
KS Naive p-value.....	0.154	0.48581
KS Statistic.....	0.11868	0.1

***** (V9) BC *****

	Before Matching	After Matching
mean treatment.....	4.9519	4.9519
mean control.....	2.9916	4.5069
std mean diff.....	39.082	8.873
mean raw eQQ diff.....	2.0337	0.72893
med raw eQQ diff.....	0.055	0.01
max raw eQQ diff.....	9.5	7.12
mean eCDF diff.....	0.089824	0.040026
med eCDF diff.....	0.066484	0.042857
max eCDF diff.....	0.23736	0.1
var ratio (Tr/Co).....	1.1009	0.91762
T-test p-value.....	0.00018579	0.4144
KS Bootstrap p-value..	< 2.22e-16	0.328
KS Naive p-value.....	7.0428e-05	0.48581
KS Statistic.....	0.23736	0.1

***** (V10) BD *****

	Before Matching	After Matching
mean treatment.....	4.52	4.52
mean control.....	2.4404	3.83
std mean diff.....	44.618	14.804
mean raw eQQ diff.....	2.0993	0.69429
med raw eQQ diff.....	0.65	0.15
max raw eQQ diff.....	8.5	6.2
mean eCDF diff.....	0.14507	0.053945
med eCDF diff.....	0.17527	0.057143
max eCDF diff.....	0.22308	0.1
var ratio (Tr/Co).....	1.4089	1.1006
T-test p-value.....	1.0928e-05	0.12636
KS Bootstrap p-value..	< 2.22e-16	0.286
KS Naive p-value.....	0.00023316	0.48581
KS Statistic.....	0.22308	0.1

***** (V11) ps *****

	Before Matching	After Matching
mean treatment.....	0.45945	0.45945
mean control.....	0.29107	0.41398

std mean diff.....	93.884	25.351
mean raw eQQ diff.....	0.16923	0.045731
med raw eQQ diff.....	0.17888	0.055288
max raw eQQ diff.....	0.2476	0.098859
mean eCDF diff.....	0.24865	0.074643
med eCDF diff.....	0.26429	0.067857
max eCDF diff.....	0.39341	0.17143
var ratio (Tr/Co).....	0.94368	1.2239
T-test p-value.....	< 2.22e-16	2.6324e-07
KS Bootstrap p-value..	< 2.22e-16	0.032
KS Naive p-value.....	1.1692e-12	0.032675
KS Statistic.....	0.39341	0.17143

***** (V12) linps *****

	Before Matching	After Matching
mean treatment.....	-0.18896	-0.18896
mean control.....	-1.0761	-0.38328
std mean diff.....	110.7	24.248
mean raw eQQ diff.....	0.89465	0.19592
med raw eQQ diff.....	0.9187	0.23821
max raw eQQ diff.....	1.5824	0.47847
mean eCDF diff.....	0.24865	0.074643
med eCDF diff.....	0.26429	0.067857
max eCDF diff.....	0.39341	0.17143
var ratio (Tr/Co).....	0.62742	1.2396
T-test p-value.....	< 2.22e-16	2.5765e-07
KS Bootstrap p-value..	< 2.22e-16	0.032
KS Naive p-value.....	1.1692e-12	0.032675
KS Statistic.....	0.39341	0.17143

Before Matching Minimum p.value: < 2.22e-16

Variable Name(s): covC covE BC BD ps linps Number(s): 3 5 9 10 11 12

After Matching Minimum p.value: 2.5765e-07

Variable Name(s): linps Number(s): 12

The `cobalt` package has some promising tools for taking this sort of output and turning it into something useful. We'll look at that approach soon. For now, some old-school stuff...

Extracting, Tabulating Standardized Differences (without `cobalt`)

We'll start by naming the covariates that the `MatchBalance` output contains...

```
covnames <- c("covA", "covB", "covC", "covD", "covE",
              "covF - Middle", "covF - High",
              "A^2", "B*C", "B*D", "raw PS", "linear PS")
```

The next step is to extract the standardized differences (using the pooled denominator to estimate, rather than the treatment-only denominator used in the main output above.)

```
pre.szd <- NULL; post.szd <- NULL
for(i in 1:length(covnames)) {
  pre.szd[i] <- mb1$BeforeMatching[[i]]$sdiff.pooled
  post.szd[i] <- mb1$AfterMatching[[i]]$sdiff.pooled
}
```

Now, we can build a table of the standardized differences:


```
match_szd <- data.frame(covnames, pre.szd, post.szd, row.names=covnames)
print(match_szd, digits=3)
```

	covnames	pre.szd	post.szd
covA	covA	14.33	7.45
covB	covB	45.44	12.82
covC	covC	-49.56	-8.44
covD	covD	23.85	-2.30
covE	covE	-48.61	-9.81
covF - Middle	covF - Middle	1.81	-11.70
covF - High	covF - High	40.47	20.99
A^2	A^2	16.94	7.89
B*C	B*C	40.01	8.87
B*D	B*D	48.26	14.80
raw PS	raw PS	92.51	25.35
linear PS	linear PS	97.21	24.25

And then, we could plot these, or their absolute values. Here's what that looks like.

A Love Plot describing Standardized Differences Before/After Matching (without cobalt)

```
ggplot(match_szd, aes(x = pre.szd, y = reorder(covnames, pre.szd))) +
  geom_point(col = "black", size = 3, pch = 1) +
  geom_point(aes(x = post.szd, y = reorder(covnames, pre.szd)),
    size = 3, col = "blue") +
  theme_bw() +
  geom_vline(aes(xintercept = 0)) +
  geom_vline(aes(xintercept = 10), linetype = "dashed", col = "red") +
  geom_vline(aes(xintercept = -10), linetype = "dashed", col = "red") +
  labs(x = "Standardized Difference (%)", y = "")
```



Using cobalt to build a “Love Plot” after Matching

```
b <- bal.tab(match1, treated ~ covA + covB + covC + covD + covE + covF +
             Asqr + BC + BD + ps + linps, data=toy, un = TRUE)
```

b

Balance Measures

	Type	Diff.Un	Diff.Adj
covA	Contin.	0.1405	0.0745
covB	Binary	0.2181	0.0643
covC	Contin.	-0.5190	-0.0844
covD	Contin.	0.2460	-0.0230
covE	Contin.	-0.5383	-0.0981
covF_1-Low	Binary	-0.1824	-0.0429
covF_2-Middle	Binary	0.0088	-0.0571
covF_3-High	Binary	0.1736	0.1000
Asqr	Contin.	0.1605	0.0789
BC	Contin.	0.3908	0.0887
BD	Contin.	0.4462	0.1480
ps	Contin.	0.9388	0.2535
linps	Contin.	1.1070	0.2425


Sample sizes

	Control	Treated
All	260	140

Matched	140	140
Unmatched	120	0

Building a Plot of Standardized Differences, with cobalt


```
p <- love.plot(b, threshold = .1, size = 1.5,
               var.order = "unadjusted",
               title = "Standardized Differences and 1:1 Matching")
p + theme_bw()
```



Building a Plot of Variance Ratios, with cobalt

```
p <- love.plot(b, stat = "v",
               threshold = 1.25, size = 1.5,
               var.order = "unadjusted",
               title = "Variance Ratios and 1:1 Matching")
p + theme_bw()
```

Variance Ratios and 1:1 Matching



Extracting, Tabulating Variance Ratios (without cobalt)

Next, we extract the variance ratios, and build a table.

```
pre.vratio <- NULL; post.vratio <- NULL
for(i in 1:length(covnames)) {
  pre.vratio[i] <- mb1$BeforeMatching[[i]]$var.ratio
  post.vratio[i] <- mb1$AfterMatching[[i]]$var.ratio
}

## Table of Variance Ratios
match_vrat <- data.frame(names = covnames, pre.vratio, post.vratio, row.names=covnames)
print(match_vrat, digits=2)
```

	names	pre.vratio	post.vratio
covA	covA	1.08	1.00
covB	covB	1.20	1.01
covC	covC	0.84	0.87
covD	covD	0.89	1.09
covE	covE	0.69	1.11
covF - Middle	covF - Middle	1.01	0.96
covF - High	covF - High	1.61	1.23
A ²	A ²	1.26	1.17
B*C	B*C	1.10	0.92
B*D	B*D	1.41	1.10
raw PS	raw PS	0.94	1.22

linear PS	linear PS	0.63	1.24
-----------	-----------	------	------

Creating a New Data Frame, Containing the Matched Sample (without cobalt)

Now, we build a new matched sample data frame in order to do some of the analyses to come. This will contain only the 280 matched subjects (140 treated and 140 control).

```
matches <- factor(rep(match1$index.treated, 2))
toy.matchedsample <- cbind(matches, toy[c(match1$index.control, match1$index.treated),])
```

Some sanity checks:

```
toy.matchedsample %>% count(treated_f)
```

```
# A tibble: 2 x 2
  treated_f     n
  <fct>      <int>
1 Treated    140
2 Control    140
```

```
head(toy.matchedsample)
```

	matches	subject	treated	covA	covB	covC	covD	covE	covF	out1.cost
1	2	T_260	0	3.08	1	10.30	9.4	10	1-Low	42
2	5	T_138	0	3.84	0	9.82	9.0	10	1-Low	58
3	11	T_190	0	2.86	0	7.50	12.0	5	3-High	39
4	14	T_235	0	3.87	1	10.20	9.5	7	2-Middle	51
5	15	T_297	0	4.01	0	9.00	12.7	13	2-Middle	49
6	17	T_261	0	5.35	1	5.56	10.3	10	2-Middle	82

	out2.event	out3.time	treated_f	covB_f	out2_f	out2	covF.Low	covF.Middle
1	No	127	Control	Has B	No Event	0	1	0
2	Yes	92	Control	No B	Event	1	1	0
3	No	105	Control	No B	No Event	0	0	0
4	No	108	Control	Has B	No Event	0	0	1
5	No	111	Control	No B	No Event	0	0	1
6	Yes	114	Control	Has B	Event	1	0	1

	covF.High	Asqr	BC	BD	ps	linps	exposure
1	0	9.4864	10.30	9.4	0.4351701	-0.2607876	0
2	0	14.7456	0.00	0.0	0.2450288	-1.1253040	0
3	1	8.1796	0.00	0.0	0.7704718	1.2109770	0
4	0	14.9769	10.20	9.5	0.6434764	0.5904848	0
5	0	16.0801	0.00	0.0	0.2989950	-0.8520883	0
6	0	28.6225	5.56	10.3	0.7069386	0.8805615	0

Rubin's Rules to Check Balance After Matching

Rubin's Rule 1

Rubin's Rule 1 states that the absolute value of the standardized difference of the linear propensity score, comparing the treated group to the control group, should be close to 0, ideally below 10%, and in any case less than 50%. If so, we may move on to Rule 2.

Recall that our result without propensity matching (or any other adjustment) was

```
rubin1.unadj
```

```
[1] 85.85784
```

To run this for our matched sample, we use:

```
rubin1.match <- with(toy.matchedsample,
  abs(100*(mean(linps[treated==1]) - mean(linps[treated==0]))/sd(linps)))
rubin1.match
```

```
[1] 25.35097
```

Here, we've at least got this value down below 50%, so we would pass Rule 1, although perhaps a different propensity score adjustment (perhaps by weighting or subclassification, or using a different matching approach) might improve this result by getting it closer to 0.

Rubin's Rule 2

Rubin's Rule 2 states that the ratio of the variance of the linear propensity score in the treated group to the variance of the linear propensity score in the control group should be close to 1, ideally between 4/5 and 5/4, but certainly not very close to or exceeding 1/2 and 2. If so, we may move on to Rule 3.

Recall that our result without propensity matching (or any other adjustment) was

```
rubin2.unadj
```

```
[1] 0.6274233
```

To run this for our matched sample, we use:

```
rubin2.match <- with(toy.matchedsample, var(linps[treated==1])/var(linps[treated==0]))
rubin2.match
```

```
[1] 1.239624
```

This is moderately promising - a substantial improvement over our unadjusted result, and now, just barely within our desired range of 4/5 to 5/4, and clearly within 1/2 to 2.

We pass Rule 2, as well.

Rubin's Rule 3

For Rubin's Rule 3, we begin by calculating regression residuals for each covariate of interest (usually, each of those included in the propensity model) regressed on a single predictor - the linear propensity score. We then look to see if the ratio of the variance of the residuals of this model for the treatment group divided by the variance of the residuals of this model for the control group is close to 1. Again, ideally this will fall between 4/5 and 5/4 for each covariate, but certainly between 1/2 and 2. If so, then the use of regression models seems well justified.

Recall that our result without propensity matching (or any other adjustment) was

```
rubin3.unadj
```

```
# A tibble: 10 x 2
  name      resid.var.ratio
  <chr>          <dbl>
1 covA          1.08
2 covB          1.49
3 covC          0.971
```

```

4 covD                1.00
5 covE                0.717
6 covF.Middle         1.03
7 covF.High           1.48
8 Asqr                1.25
9 BC                  1.32
10 BD                 1.76

```

After propensity matching, we use this code to assess Rubin's 3rd Rule in our matched sample.

```

cov.sub <- dplyr::select(toy.matchedsample,
                        covA, covB, covC, covD, covE,
                        covF.Middle, covF.High, Asqr, BC, BD)

toy.matchedsample$exposure <- toy.matchedsample$treated

rubin3.matched <- rubin3(data = toy.matchedsample, covlist = cov.sub, linps = linps)

rubin3.matched

```

```

# A tibble: 10 x 2
  name      resid.var.ratio
  <chr>          <dbl>
1 covA          1
2 covB          1.15
3 covC          0.793
4 covD          1.11
5 covE          0.922
6 covF.Middle   0.933
7 covF.High     1.24
8 Asqr          1.17
9 BC            0.982
10 BD           1.28

```

It looks like the results are basically unchanged, except that `covF.High` is improved. The dotplot of these results comparing pre- to post-matching is shown below.


A Cleveland Dot Chart of the Rubin's Rule 3 Results Pre vs. Post-Match

```

rubin3.both <- bind_rows(rubin3.unadj, rubin3.matched)
rubin3.both$source <- c(rep("Unmatched",10), rep("Matched", 10))

ggplot(rubin3.both, aes(x = resid.var.ratio, y = name, col = source)) +
  geom_point(size = 2) +
  theme_bw() +
  xlim(0.5, 2.0) +
  geom_vline(aes(xintercept = 1)) +
  geom_vline(aes(xintercept = 4/5), linetype = "dashed", col = "red") +
  geom_vline(aes(xintercept = 5/4), linetype = "dashed", col = "red") +
  labs(x = "Residual Variance Ratio", y = "")

```



Some improvement to report, overall.

Task 5. After matching, estimate the causal effect of treatment on ...

Outcome 1 (a continuous outcome)

Approach 1. Automated Approach from the Matching package - ATT Estimate

First, we'll look at the essentially automatic answer which can be obtained when using the `Matching` package and inserting an outcome `Y`. For a continuous outcome, this is often a reasonable approach.

```
X <- toy$linps ## matching on the linear propensity score
Tr <- as.logical(toy$treated)
Y <- toy$out1.cost
match1.out1 <- Match(Y=Y, Tr=Tr, X=X, M = 1, replace=FALSE, ties=FALSE)
summary(match1.out1)
```

```
Estimate... 9.8071
SE..... 1.6111
T-stat..... 6.0873
p.val..... 1.1484e-09
```

```
Original number of observations..... 400
```

```
Original number of treated obs..... 140
Matched number of observations..... 140
Matched number of observations (unweighted). 140
```

The estimate is 9.81 with standard error 1.61. We can obtain an approximate 95% confidence interval by adding and subtracting 1.96 times (or just double) the standard error (SE) to the point estimate, 9.81. Here, using the 1.96 figure, that would yields an approximate 95% CI of (6.65, 12.96).

Approach 2. Automated Approach from the Matching package - ATE Estimate

```
match1.out1.ATE <- Match(Y=Y, Tr=Tr, X=X, M = 1, replace=FALSE, ties=FALSE, estimand="ATE")
summary(match1.out1.ATE)
```

```
Estimate... 9.8393
SE..... 1.1568
T-stat..... 8.5053
p.val..... < 2.22e-16
```

```
Original number of observations..... 400
Original number of treated obs..... 140
Matched number of observations..... 280
Matched number of observations (unweighted). 280
```

And our 95% CI for this ATE estimate would be $9.84 \pm 1.96(1.16)$, or (7.57, 12.11), but we'll stick with the ATT estimate for now.

ATT vs. ATE: Definitions

- Informally, the **average treatment effect on the treated** (ATT) estimate describes the difference in potential outcomes (between treated and untreated subjects) summarized across the population of people who actually received the treatment.
 - In our initial match, we identified a unique and nicely matched control patient for each of the 140 people in the treated group. We have a 1:1 match on the treated, and thus can describe subjects across that set of treated patients reasonably well.
- On the other hand the **average treatment effect** (ATE) refers to the difference in potential outcomes summarized across the entire population, including those who did not receive the treatment.
 - In our ATE match, we have less success, in part because if we match to the treated patients in a 1:1 way, we'll have an additional 120 unmatched control patients, about whom we can describe results only vaguely. We could consider matching up control patients to treated patients, perhaps combined with a willingness to re-use some of the treated patients to get a better estimate across the whole population.

Approach 3. Mirroring the Paired T test in a Regression Model

We can mirror the paired t test result in a regression model that treats the match identifier as a fixed factor in a linear model, as follows. This takes the pairing into account, but treating pairing as a fixed, rather than random, factor, isn't really satisfactory as a solution, although it does match the paired t test.

```
adj.m.out1 <- lm(out1.cost ~ treated + factor(matches), data=toy.matchedsample)

adj.m.out1.tidy <- tidy(adj.m.out1, conf.int = TRUE) %>%
  filter(term == "treated")
```

```
adj.m.out1.tidy
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic    p.value conf.low conf.high
<chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
1 treated      9.72      1.62      6.00 0.0000000158      6.52      12.9
```

So, this regression approach produces an estimate that is exactly the same as the paired t test², but this isn't something I'm completely comfortable with.

Approach 4. A Mixed Model to account for 1:1 Matching

What I think of as a more appropriate result comes from a mixed model where the matches are treated as a random factor, but the treatment group is treated as a fixed factor. This is developed like this, using the `lme4` package. Note that we have to create a factor variable to represent the matches, since that's the only thing that `lme4` understands.

```
toy.matchedsample$matches.f <- as.factor(toy.matchedsample$matches)
## Need to use matches as a factor in R here

matched_mixedmodel.out1 <- lmer(out1.cost ~ treated + (1 | matches.f), data=toy.matchedsample)
summary(matched_mixedmodel.out1); confint(matched_mixedmodel.out1)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: out1.cost ~ treated + (1 | matches.f)
Data: toy.matchedsample
```

```
REML criterion at convergence: 2296.1
```

```
Scaled residuals:
    Min       1Q   Median       3Q      Max
-2.43885 -0.69419 -0.01592  0.63684  2.17289
```

```
Random effects:
 Groups   Name      Variance Std.Dev.
 matches.f (Intercept) 38.17    6.178
 Residual              183.46   13.545
Number of obs: 280, groups: matches.f, 140
```

```
Fixed effects:
              Estimate Std. Error t value
(Intercept)   46.921     1.258   37.292
treated         9.721     1.619    6.005
```

```
Correlation of Fixed Effects:
      (Intr)
treated -0.643
```

```
              2.5 %    97.5 %
.sig01       2.429038  8.837784
.sigma       12.057491 15.245694
(Intercept)  44.455490 49.387367
treated       6.537961 12.904896
```

²I'll leave checking that this is true as an exercise for the curious.

The tidy approach works with this linear mixed model, so we have:

```
res_matched_1 <- tidy(matched_mixedmodel.out1, conf.int = TRUE, conf.level = 0.95) %>%
  filter(term == "treated")
```

```
res_matched_1
```

```
# A tibble: 1 x 7
```

```
  term      estimate std.error statistic conf.low conf.high group
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl> <chr>
1 treated    9.72      1.62      6.00    6.55   12.9 fixed
```

Our estimate is 9.72, with 95% CI ranging from 6.55 to 12.89.

Practically, does any of this matter in this example?

Not much in this example, no, as long as you stick to ATT approaches.

	Approach	Effect Estimate	Standard Error	95% CI
	“Automated” ATT via Match	9.81	1.61	(6.65, 12.96)
	Linear Model (pairs as fixed factor)	9.72	1.62	(6.52, 12.92)
	Mixed Model (pairs as random factor)	9.72	1.62	(6.55, 12.89)

Outcome 2 (a binary outcome)

Approach 1. Automated Approach from the Matching package (ATT)

First, we’ll look at the essentially automatic answer which can be obtained when using the **Matching** package and inserting an outcome Y. For a binary outcome, this is often a reasonable approach, especially if you don’t wish to adjust for any other covariate, and the result will be expressed as a risk difference, rather than as a relative risk or odds ratio. Note that I have used the 0-1 version of Outcome 2, rather than a factor version. The estimate produced is the difference in risk associated with `out2 = 1` (Treated subjects) minus `out2 = 0` (Controls.)

```
X <- toy$linps ## matching on the linear propensity score
Tr <- as.logical(toy$treated)
Y <- toy$out2
match1_out2 <- Match(Y=Y, Tr=Tr, X=X, M = 1, replace=FALSE, ties=FALSE)
summary(match1_out2)
```

```
Estimate... 0.14286
SE..... 0.061918
T-stat..... 2.3072
p.val..... 0.021043
```

```
Original number of observations..... 400
Original number of treated obs..... 140
Matched number of observations..... 140
Matched number of observations (unweighted). 140
```

As in the continuous case, we obtain an approximate 95% confidence interval by adding and subtracting 1.96 times (or just double) the standard error (SE) to the point estimate. The estimated effect on the risk difference is 0.143 with standard error 0.062 and 95% CI (0.021, 0.264).

Approach 2. Using the matched sample to perform a conditional logistic regression

Since we have the matched sample available, we can simply perform a conditional logistic regression to estimate the treatment effect in terms of a log odds ratio (or, by exponentiating, an odds ratio.) Again, I use the 0/1 version of both the outcome and treatment indicator. The key modeling function `clogit` is part of the `survival` package.

```
adj.m.out2 <- clogit(out2 ~ treated + strata(matches), data=toy.matchedsample)
summary(adj.m.out2)
```

Call:

```
coxph(formula = Surv(rep(1, 280L), out2) ~ treated + strata(matches),
      data = toy.matchedsample, method = "exact")
```

```
n= 280, number of events= 145
```

```
      coef exp(coef) se(coef)      z Pr(>|z|)
treated 0.5039    1.6552  0.2352 2.143  0.0322 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
      exp(coef) exp(-coef) lower .95 upper .95
treated    1.655    0.6042    1.044    2.624
```

```
Concordance= 0.623 (se = 0.078 )
Rsquare= 0.017 (max possible= 0.317 )
Likelihood ratio test= 4.74 on 1 df,  p=0.03
Wald test               = 4.59 on 1 df,  p=0.03
Score (logrank) test = 4.69 on 1 df,  p=0.03
```

The odds ratio in the `exp(coef)` section above is the average causal effect estimate - it describes the odds of having an event (out2) occur associated with being a treated subject, as compared to the odds of the event when a control subject.

I tidied this, as follows, without `conf.int = TRUE`, and got ...

```
adj.m.out2_tidy <- tidy(adj.m.out2, exponentiate = TRUE)
```

```
adj.m.out2_tidy
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated    1.66      0.235      2.14  0.0322    1.04    2.62
```

Our point estimate is 1.66, with standard error 0.24, and 95% CI ranging from 1.04 to 2.62.

- I'll use this conditional logistic regression approach to summarize the findings with regard to an odds ratio in my summary of matching results to come.

Outcome 3 (a time-to-event outcome)

Approach 1. Automated Approach from the Matching package

Again, we'll start by thinking about the essentially automatic answer which can be obtained when using the `Match` function. The problem here is that this approach doesn't take into account the right censoring at all,

and assumes that all of the specified times in Outcome 3 are observed. This causes the result (or the ATE version) to be non-sensical, given what we know about the data. So I don't recommend you use this approach when dealing with a time-to-event outcome.

And as a result, I won't even show it here.

Approach 2. A stratified Cox proportional hazards model

Since we have the matched sample, we can use a stratified Cox proportional hazards model to compare the treatment groups on our time-to-event outcome, while accounting for the matched pairs. The main results will be a relative hazard rate estimate, with 95% CI. Again, I use the 0/1 numeric version of the event indicator (out2), and of the treatment indicator (treated) here.

```
adj.m.out3 <- coxph(Surv(out3.time, out2) ~ treated + strata(matches), data=toy.matchedsample)
summary(adj.m.out3)
```

Call:

```
coxph(formula = Surv(out3.time, out2) ~ treated + strata(matches),
      data = toy.matchedsample)
```

```
n= 280, number of events= 145
```

```
      coef exp(coef) se(coef)      z Pr(>|z|)
treated 0.5845    1.7941   0.2140 2.731 0.00631 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
      exp(coef) exp(-coef) lower .95 upper .95
treated    1.794    0.5574    1.179    2.729
```

```
Concordance= 0.642 (se = 0.07 )
Rsquare= 0.027 (max possible= 0.375 )
Likelihood ratio test= 7.78 on 1 df,  p=0.005
Wald test               = 7.46 on 1 df,  p=0.006
Score (logrank) test = 7.67 on 1 df,  p=0.006
```

I tidied this with ...

```
adj.m.out3_tidy <- tidy(adj.m.out3, exponentiate = TRUE)
```

```
adj.m.out3_tidy
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated    1.79      0.214      2.73 0.00631    1.18    2.73
```


Our point estimate for the relative hazard rate (from the `exp(coef)` section of the summary output) is 1.79, with standard error 0.21, and 95% CI ranging from 1.18 to 2.73.

Checking the proportional hazards assumption looks all right.

```
cox.zph(adj.m.out3) # Quick check for proportional hazards assumption
```

```
      rho chisq      p
treated -0.0421 0.257 0.612
```

```
plot(cox.zph(adj.m.out3), var="treated")
```



Results So Far (After Propensity Matching)

So, here's our summary again, now incorporating both our unadjusted results and the results after matching. Automated results and my favorite of our various non-automated approaches are shown. Note that I've left out the "automated" approach for a time-to-event outcome entirely, so as to discourage you from using it.

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Relative Hazard Rate)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)
After 1:1 PS Match (Match: Automated)	9.81 (6.65, 12.96)	0.143 (0.021, 0.264)	N/A	N/A
After 1:1 PS Match ("Regression" Models)	9.72 (6.55, 12.89)	N/A	1.66 (1.04, 2.62)	1.79 (1.18, 2.73)

Task 6. Subclassify by PS quintile, then display post-subclassification balance

First, we divide the data by the propensity score into 5 strata of equal size using the `cut2` function from the `Hmisc` package. Then we create a `quintile` variable which specifies 1 = lowest propensity scores to 5 = highest.

```
toy$stratum <- Hmisc::cut2(toy$ps, g=5)
```

```
toy %>% group_by(stratum) %>% skim(ps) ## sanity check
```

Skim summary statistics

n obs: 400

n variables: 25

group variables: stratum

```
-- Variable type:numeric -----
      stratum variable missing complete  n mean   sd   p0   p25  p50
[0.0345,0.170)      ps        0       80 80 0.1  0.036 0.034 0.074 0.1
[0.1698,0.259)      ps        0       80 80 0.21 0.025 0.17  0.2   0.22
[0.2588,0.386)      ps        0       80 80 0.31 0.038 0.26  0.29  0.31
[0.3861,0.545)      ps        0       80 80 0.46 0.045 0.39  0.43  0.46
[0.5453,0.834]      ps        0       80 80 0.66 0.067 0.55  0.6   0.65
  p75 p100
0.13 0.17
0.24 0.26
0.35 0.38
0.49 0.54
0.71 0.83
```

```
toy$quintile <- factor(toy$stratum, labels=1:5)
```

```
toy %>% count(stratum, quintile) ## sanity check
```

A tibble: 5 x 3

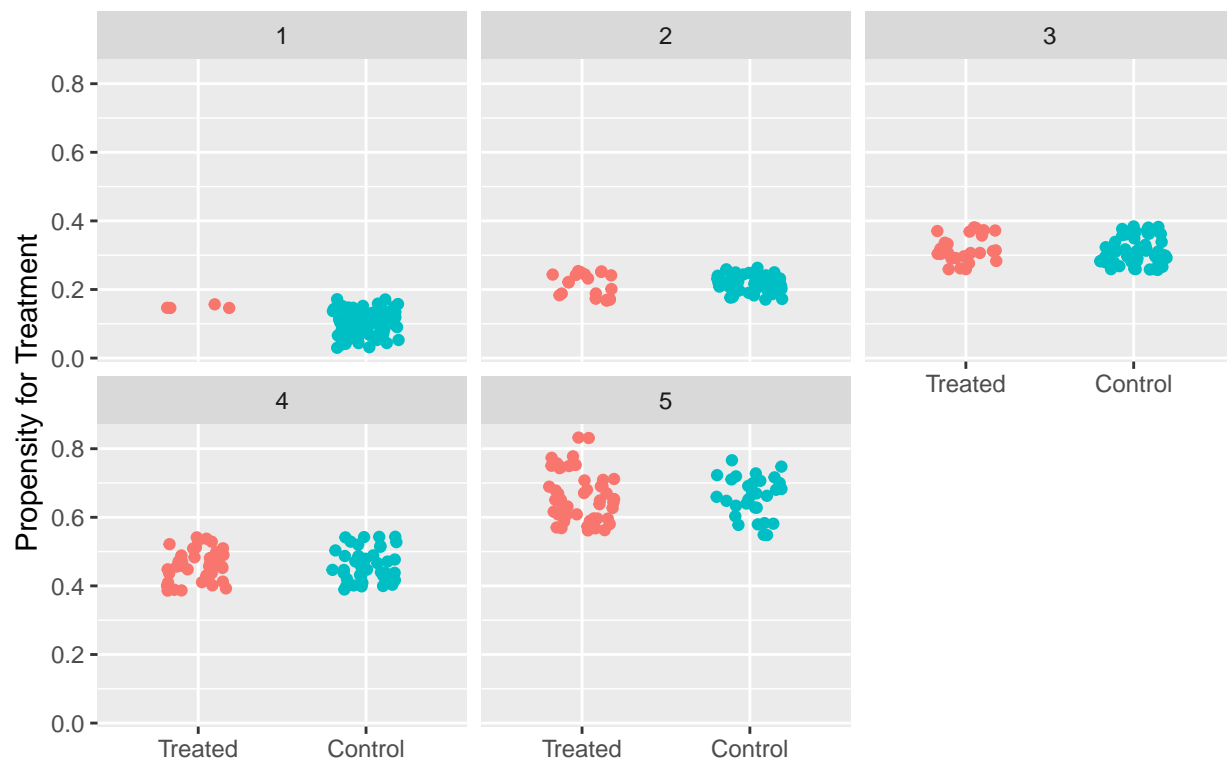
	stratum	quintile	n
	<fct>	<fct>	<int>
1	[0.0345,0.170)	1	80
2	[0.1698,0.259)	2	80
3	[0.2588,0.386)	3	80
4	[0.3861,0.545)	4	80
5	[0.5453,0.834]	5	80

Check Balance and Propensity Score Overlap in Each Quintile

We want to check the balance and propensity score overlap for each stratum (quintile.) I'll start with a set of faceted, jittered plots to look at overlap.

```
ggplot(toy, aes(x = treated_f, y = round(ps,2), group = quintile, color = treated_f)) +
  geom_jitter(width = 0.2) +
  guides(color = FALSE) +
  facet_wrap(~ quintile) +
  labs(x = "", y = "Propensity for Treatment",
       title = "Quintile Subclassification in the Toy Example")
```

Quintile Subclassification in the Toy Example



It can be helpful to know how many observations (by exposure group) are in each quintile.

```
toy %>% count(quintile, treated_f)
```

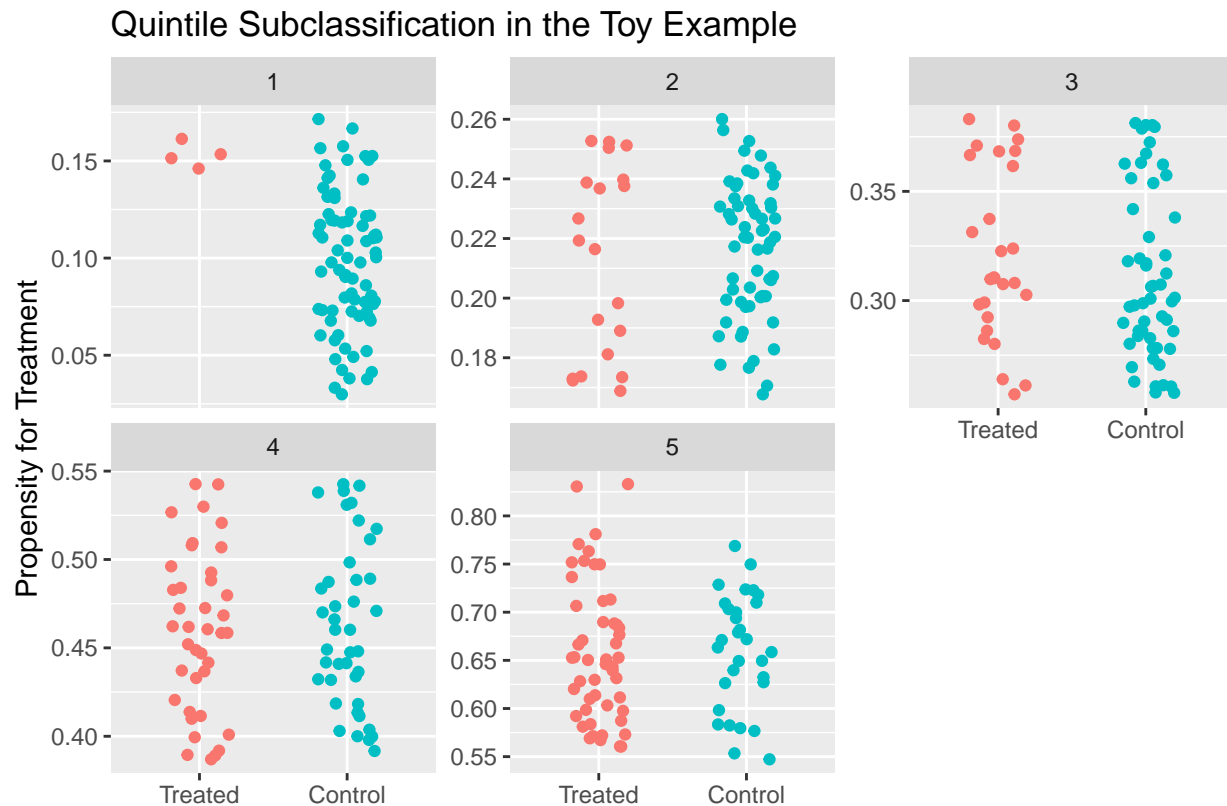
```
# A tibble: 10 x 3
  quintile treated_f    n
  <fct>      <fct>    <int>
1 1         Treated      4
2 1         Control     76
3 2         Treated     20
4 2         Control     60
5 3         Treated     27
6 3         Control     53
7 4         Treated     39
8 4         Control     41
9 5         Treated     50
10 5        Control     30
```

With only 4 “treated” subjects in Quintile 1, I am concerned that we won’t be able to do much there to create balance.

The overlap may show a little better in the plot if you free up the y axes...

```
ggplot(toy, aes(x = treated_f, y = round(ps,2), group = quintile, color = treated_f)) +
  geom_jitter(width = 0.2) +
  guides(color = FALSE) +
  facet_wrap(~ quintile, scales = "free_y") +
  labs(x = "", y = "Propensity for Treatment",
```

```
title = "Quintile Subclassification in the Toy Example")
```



Creating a Standardized Difference Calculation Function

We'll need to be able to calculate standardized differences in this situation so I've created a simple `szd` function to do this - using the average denominator method.

```
szd <- function(covlist, g) {
  covlist2 <- as.matrix(covlist)
  g <- as.factor(g)
  res <- NA
  for(i in 1:ncol(covlist2)) {
    cov <- as.numeric(covlist2[,i])
    num <- 100*diff(tapply(cov, g, mean, na.rm=TRUE))
    den <- sqrt(mean(tapply(cov, g, var, na.rm=TRUE)))
    res[i] <- round(num/den,2)
  }
  names(res) <- names(covlist)
  res
}
```

Creating the Five Subsamples, by PS Quintile

Next, we split the complete sample into the five quintiles.

```
## Divide the sample into the five quintiles
quin1 <- filter(toy, quintile==1)
quin2 <- filter(toy, quintile==2)
quin3 <- filter(toy, quintile==3)
quin4 <- filter(toy, quintile==4)
quin5 <- filter(toy, quintile==5)
```

Standardized Differences in Each Quintile, and Overall

Now, we'll calculate the standardized differences for each covariate (note that we're picking up two of the indicators for our multi-categorical covF) within each quintile, as well as overall.

```
covs <- c("covA", "covB", "covC", "covD", "covE", "covF.Middle",
          "covF.High", "Asqr", "BC", "BD", "ps", "linps")
d.q1 <- szd(quin1[covs], quin1$treated)
d.q2 <- szd(quin2[covs], quin2$treated)
d.q3 <- szd(quin3[covs], quin3$treated)
d.q4 <- szd(quin4[covs], quin4$treated)
d.q5 <- szd(quin5[covs], quin5$treated)
d.all <- szd(toy[covs], toy$treated)

toy.szd <- data_frame(covs, Overall = d.all, Q1 = d.q1, Q2 = d.q2, Q3 = d.q3, Q4 = d.q4, Q5 = d.q5)
toy.szd <- gather(toy.szd, "quint", "sz.diff", 2:7)
toy.szd
```

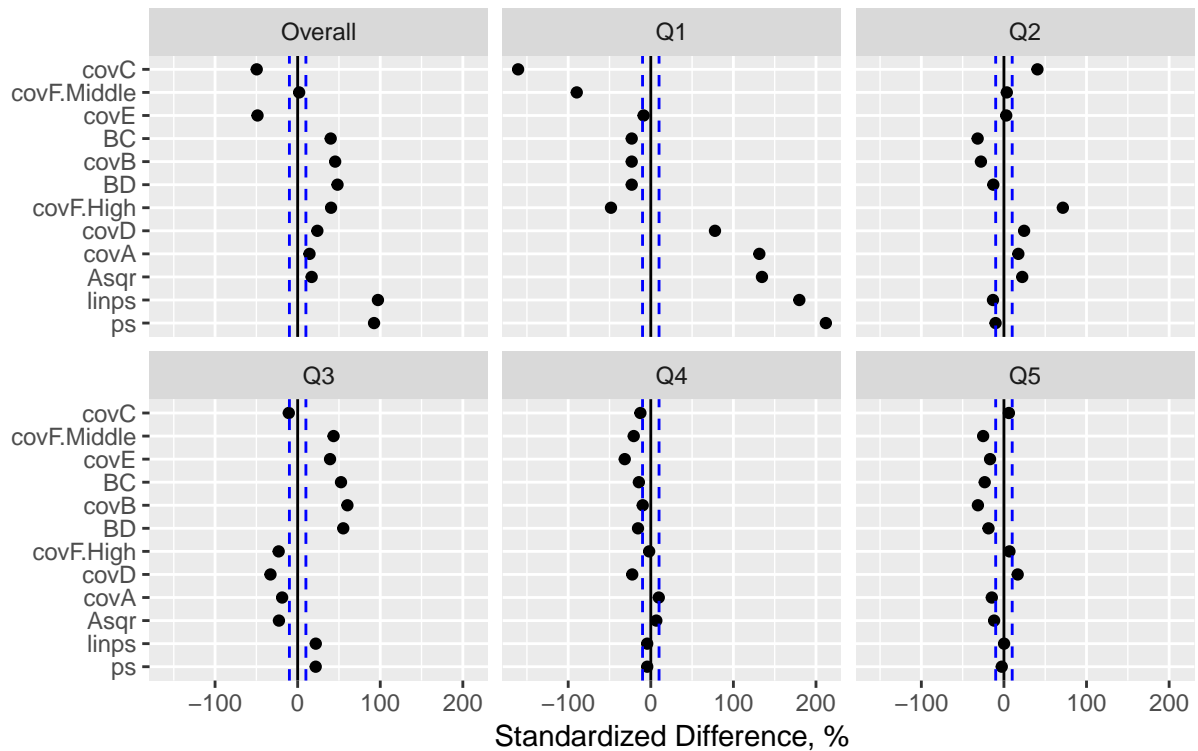
```
# A tibble: 72 x 3
  covs      quint  sz.diff
<chr>    <chr>    <dbl>
1 covA      Overall    14.3
2 covB      Overall    45.4
3 covC      Overall   -49.6
4 covD      Overall    23.8
5 covE      Overall   -48.6
6 covF.Middle Overall     1.81
7 covF.High Overall    40.5
8 Asqr      Overall    16.9
9 BC        Overall    40.0
10 BD       Overall    48.3
# ... with 62 more rows
```

Plotting the Standardized Differences

```
ggplot(toy.szd, aes(x = sz.diff, y = reorder(covs, -sz.diff), group = quint)) +
  geom_point() +
  geom_vline(xintercept = 0) +
  geom_vline(xintercept = c(-10,10), linetype = "dashed", col = "blue") +
  facet_wrap(~ quint) +
  labs(x = "Standardized Difference, %", y = "",
       title = "Comparing Standardized Differences by PS Quintile",
       subtitle = "The toy example")
```


Comparing Standardized Differences by PS Quintile

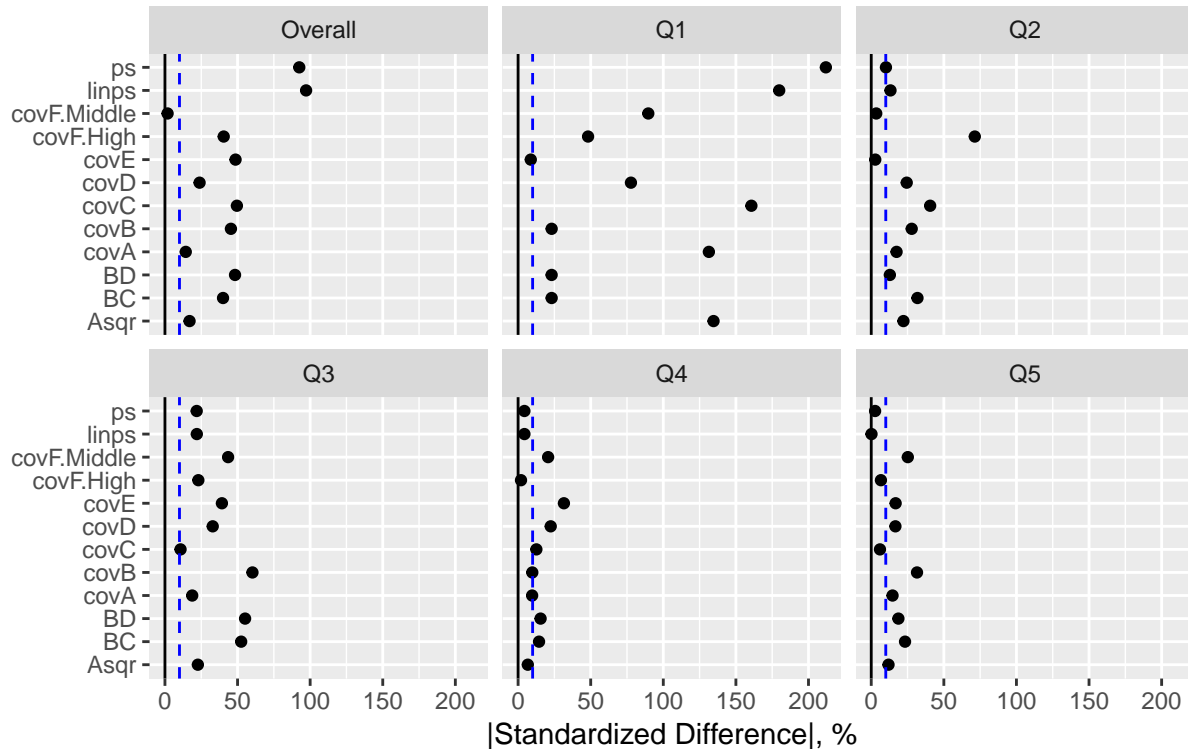
The toy example



```
ggplot(toy.szd, aes(x = abs(szd.diff), y = covs, group = quint)) +
  geom_point() +
  geom_vline(xintercept = 0) +
  geom_vline(xintercept = 10, linetype = "dashed", col = "blue") +
  facet_wrap(~ quint) +
  labs(x = "|Standardized Difference|, %", y = "",
       title = "Absolute Standardized Differences by PS Quintile",
       subtitle = "The toy example")
```

Absolute Standardized Differences by PS Quintile

The toy example



Checking Rubin's Rules Post-Subclassification

Rubin's Rule 1

As a reminder, prior to adjustment, Rubin's Rule 1 for the toy example was:

```
rubin1.unadj <- with(toy,
  abs(100*(mean(linps[treated==1]) -
    mean(linps[treated==0]))/sd(linps)))
rubin1.unadj
```

```
[1] 85.85784
```

After propensity score subclassification, we can obtain the same summary within each of the five quintiles...

```
rubin1.q1 <- with(quin1, abs(100*(mean(linps[treated==1]) -
  mean(linps[treated==0]))/sd(linps)))
rubin1.q2 <- with(quin2, abs(100*(mean(linps[treated==1]) -
  mean(linps[treated==0]))/sd(linps)))
rubin1.q3 <- with(quin3, abs(100*(mean(linps[treated==1]) -
  mean(linps[treated==0]))/sd(linps)))
rubin1.q4 <- with(quin4, abs(100*(mean(linps[treated==1]) -
  mean(linps[treated==0]))/sd(linps)))
rubin1.q5 <- with(quin5, abs(100*(mean(linps[treated==1]) -
  mean(linps[treated==0]))/sd(linps)))
```

```
rubin1.sub <- c(rubin1.q1, rubin1.q2, rubin1.q3, rubin1.q4, rubin1.q5)
names(rubin1.sub)=c("Q1", "Q2", "Q3", "Q4", "Q5")
```

```
rubin1.sub
```

	Q1	Q2	Q3	Q4	Q5
	125.831381	14.775967	22.061011	4.384187	0.176807

It was always a long shot that subclassification alone would reduce all of these values below 10%, but I had hoped to get them all below 50%. With only 4 “treated” subjects in Quintile 1, though, the task was too tough.

Rubin’s Rule 2

As a reminder, prior to adjustment, Rubin’s Rule 2 for the toy example was:

```
rubin2.unadj <- with(toy, var(linps[treated==1])/var(linps[treated==0]))
rubin2.unadj
```

```
[1] 0.6274233
```

After Subclassification, we can obtain the same summary within each of the five quintiles...

```
rubin2.q1 <- with(quin1, var(linps[treated==1])/var(linps[treated==0]))
rubin2.q2 <- with(quin2, var(linps[treated==1])/var(linps[treated==0]))
rubin2.q3 <- with(quin3, var(linps[treated==1])/var(linps[treated==0]))
rubin2.q4 <- with(quin4, var(linps[treated==1])/var(linps[treated==0]))
rubin2.q5 <- with(quin5, var(linps[treated==1])/var(linps[treated==0]))
```

```
rubin2.sub <- c(rubin2.q1, rubin2.q2, rubin2.q3, rubin2.q4, rubin2.q5)
names(rubin2.sub)=c("Q1", "Q2", "Q3", "Q4", "Q5")
```

```
rubin2.sub
```

	Q1	Q2	Q3	Q4	Q5
	0.006547378	2.170717727	1.054126217	0.925867014	1.600734926

Some of these variance ratios are actually a bit further from 1 than the full data set. Again, with a small sample size like this, subclassification looks like a weak choice. At most, three of the quintiles (3-4 and maybe 5) show OK variance ratios after propensity score subclassification.

Rubin’s Rule 3

Prior to propensity adjustment, recall that Rubin’s Rule 3 summaries were:

```
covs <- c("covA", "covB", "covC", "covD", "covE",
          "covF.Middle", "covF.High", "Asqr", "BC", "BD")
rubin3.unadj <- rubin3(data=toy, covlist=toy[covs])
```

After subclassification, then, Rubin’s Rule 3 summaries within each quintile are:

```
rubin3.q1 <- rubin3(data=quin1, covlist=quin1[covs])
rubin3.q2 <- rubin3(data=quin2, covlist=quin2[covs])
rubin3.q3 <- rubin3(data=quin3, covlist=quin3[covs])
rubin3.q4 <- rubin3(data=quin4, covlist=quin4[covs])
rubin3.q5 <- rubin3(data=quin5, covlist=quin5[covs])
```

```

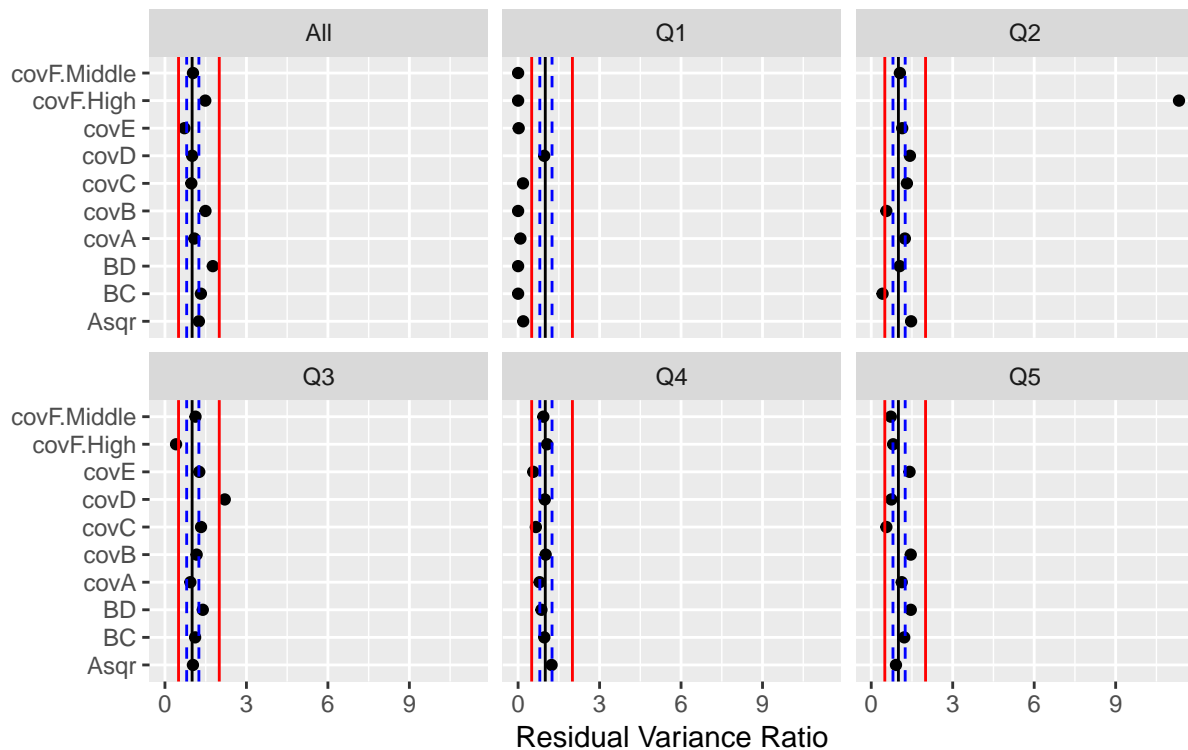
toy.rubin3 <- data_frame(covs, All = rubin3.unadj$resid.var.ratio,
                        Q1 = rubin3.q1$resid.var.ratio,
                        Q2 = rubin3.q2$resid.var.ratio,
                        Q3 = rubin3.q3$resid.var.ratio,
                        Q4 = rubin3.q4$resid.var.ratio,
                        Q5 = rubin3.q5$resid.var.ratio)

toy.rubin3 <- gather(toy.rubin3, "quint", "rubin3", 2:7)

ggplot(toy.rubin3, aes(x = rubin3, y = covs, group = quint)) +
  geom_point() +
  geom_vline(xintercept = 1) +
  geom_vline(xintercept = c(0.8, 1.25), linetype = "dashed", col = "blue") +
  geom_vline(xintercept = c(0.5, 2), col = "red") +
  facet_wrap(~ quint) +
  labs(x = "Residual Variance Ratio", y = "",
       title = "Residual Variance Ratios by PS Quintile",
       subtitle = "Rubin's Rule 3: The toy example")

```

Residual Variance Ratios by PS Quintile
Rubin's Rule 3: The toy example



Most of the residual variance ratios are in the range of (0.5, 2) in quintiles 2-5, with the exception of the covF.high indicator in Quintile 2. Quintile 1 is certainly problematic in this regard.

Task 7. After subclassifying, what is the estimated average treatment effect?

... on Outcome 1 [a continuous outcome]

First, we'll find the estimated average causal effect (and standard error) within each quintile via linear regression.

```
quin1.out1 <- lm(out1.cost ~ treated, data=quin1)
quin2.out1 <- lm(out1.cost ~ treated, data=quin2)
quin3.out1 <- lm(out1.cost ~ treated, data=quin3)
quin4.out1 <- lm(out1.cost ~ treated, data=quin4)
quin5.out1 <- lm(out1.cost ~ treated, data=quin5)

coef(summary(quin1.out1)); coef(summary(quin2.out1)); coef(summary(quin3.out1)); coef(summary(quin4.out1)); coef(summary(quin5.out1))
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	46.763158	1.283162	36.443677	9.663430e-51
treated	-4.013158	5.738477	-0.699342	4.864186e-01

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	45.5	1.445801	31.47043	4.383196e-46
treated	7.6	2.891603	2.62830	1.033042e-02

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	45.000000	1.804463	24.938163	6.523421e-39
treated	8.444444	3.106069	2.718691	8.074096e-03

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	48.097561	2.775464	17.329555	1.814301e-28
treated	9.287054	3.975103	2.336306	2.204426e-02

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	52.70	2.681145	19.655781	5.998878e-32
treated	7.62	3.391410	2.246853	2.747662e-02

Just looking at these results, it doesn't look like combining quintile 1 with the others is a good idea. I'll do it here, to show the general idea, but I'm not satisfied with the results. There is certainly a cleverer way to accomplish this using the `broom` package, or maybe a little programming with `purrr`.

Next, we find the mean of the five quintile-specific estimated regression coefficients

```
est.st <- (coef(quin1.out1)[2] + coef(quin2.out1)[2] + coef(quin3.out1)[2] +
           coef(quin4.out1)[2] + coef(quin5.out1)[2])/5
est.st
```

```
treated
5.787668
```

To get the combined standard error estimate, we do the following:

```
se.q1 <- summary(quin1.out1)$coefficients[2,2]
se.q2 <- summary(quin2.out1)$coefficients[2,2]
se.q3 <- summary(quin3.out1)$coefficients[2,2]
se.q4 <- summary(quin4.out1)$coefficients[2,2]
se.q5 <- summary(quin5.out1)$coefficients[2,2]
```

```
se.st <- sqrt((se.q1^2 + se.q2^2 + se.q3^2 + se.q4^2 + se.q5^2)*(1/25))
se.st
```

```
[1] 1.769093
```

The resulting 95% confidence Interval for the average causal treatment effect is then:

```
strat.result1 <- data_frame(estimate = est.st,
                             conf.low = est.st - 1.96*se.st,
                             conf.high = est.st + 1.96*se.st)
strat.result1
```

```
# A tibble: 1 x 3
  estimate conf.low conf.high
    <dbl>    <dbl>    <dbl>
1     5.79     2.32     9.26
```

Again, I don't trust this estimate in this setting because the balance (especially in Quintile 1) is too weak.

... on Outcome 2 [a binary outcome]

First, we find the estimated average causal effect (and standard error) within each quintile via logistic regression:

```
quin1.out2 <- glm(out2 ~ treated, data=quin1, family=binomial())
quin2.out2 <- glm(out2 ~ treated, data=quin2, family=binomial())
quin3.out2 <- glm(out2 ~ treated, data=quin3, family=binomial())
quin4.out2 <- glm(out2 ~ treated, data=quin4, family=binomial())
quin5.out2 <- glm(out2 ~ treated, data=quin5, family=binomial())
```

```
coef(summary(quin1.out2)); coef(summary(quin2.out2)); coef(summary(quin3.out2)); coef(summary(quin4.out2)); coef(summary(quin5.out2))
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.8347977	0.2496921	-3.3433088	0.0008278571
treated	0.8347977	1.0307018	0.8099314	0.4179796183

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.3364722	0.2618615	-1.284925	0.1988186
treated	1.1837701	0.5537747	2.137638	0.0325461

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.1892420	0.2759519	-0.685779	0.49285245
treated	0.8823892	0.4927637	1.790694	0.07334233

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.3448405	0.3170019	-1.087818	0.2766753
treated	0.6026696	0.4525133	1.331827	0.1829169

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.2682640	0.3684322	0.7281230	0.4665383
treated	-0.1882213	0.4646186	-0.4051092	0.6853973

Next, we find the mean of the five quintile-specific estimated logistic regression coefficients

```
est.st <- (coef(quin1.out2)[2] + coef(quin2.out2)[2] + coef(quin3.out2)[2] +
           coef(quin4.out2)[2] + coef(quin5.out2)[2])/5
est.st ## this is the estimated log odds ratio
```

```
treated
0.6630811
```

```
## And we exponentiate this to get the overall odds ratio estimate
exp(est.st)
```

```
treated
1.940763
```

To get the combined standard error estimate across the five quintiles, we do the following:

```
se.q1 <- summary(quin1.out2)$coefficients[2,2]
se.q2 <- summary(quin2.out2)$coefficients[2,2]
se.q3 <- summary(quin3.out2)$coefficients[2,2]
se.q4 <- summary(quin4.out2)$coefficients[2,2]
se.q5 <- summary(quin5.out2)$coefficients[2,2]
se.st <- sqrt((se.q1^2 + se.q2^2 + se.q3^2 + se.q4^2 + se.q5^2)*(1/25))
se.st
```

```
[1] 0.2851293
```

```
## Of course, this standard error is also on the log odds ratio scale
```

Now, we obtain a 95% Confidence Interval for the Average Causal Effect of our treatment (as an Odds Ratio) through combination and exponentiation, as follows:

```
strat.result2 <- data_frame(estimate = exp(est.st),
                           conf.low = exp(est.st - 1.96*se.st),
                           conf.high = exp(est.st + 1.96*se.st))
strat.result2
```

```
# A tibble: 1 x 3
  estimate conf.low conf.high
    <dbl>    <dbl>    <dbl>
1    1.94    1.11    3.39
```

... on Outcome 3 [a time to event]

Subjects with `out2.event` = “Yes” are truly observed events, while those with `out2.event` == “No” are censored before an event can happen to them.

The Cox model comparing treated to control, stratifying on quintile, is...

```
adj.s.out3 <- coxph(Surv(out3.time, out2) ~ treated + strata(quintile), data=toy)
summary(adj.s.out3) ## exp(coef) gives relative hazard associated with treatment
```

Call:

```
coxph(formula = Surv(out3.time, out2) ~ treated + strata(quintile),
      data = toy)
```

```
n= 400, number of events= 188
```

```
      coef exp(coef) se(coef)      z Pr(>|z|)
treated 0.6817    1.9772   0.1718 3.968 7.25e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
exp(coef) exp(-coef) lower .95 upper .95
```

```

treated      1.977      0.5058      1.412      2.769

Concordance= 0.582 (se = 0.019 )
Rsquare= 0.038 (max possible= 0.97 )
Likelihood ratio test= 15.66 on 1 df, p=8e-05
Wald test           = 15.74 on 1 df, p=7e-05
Score (logrank) test = 16.13 on 1 df, p=6e-05

strat.result3 <- tidy(adj.s.out3, exponentiate = TRUE)

```

Checking the Proportional Hazards Assumption

The proportional hazards assumption may be problematic.

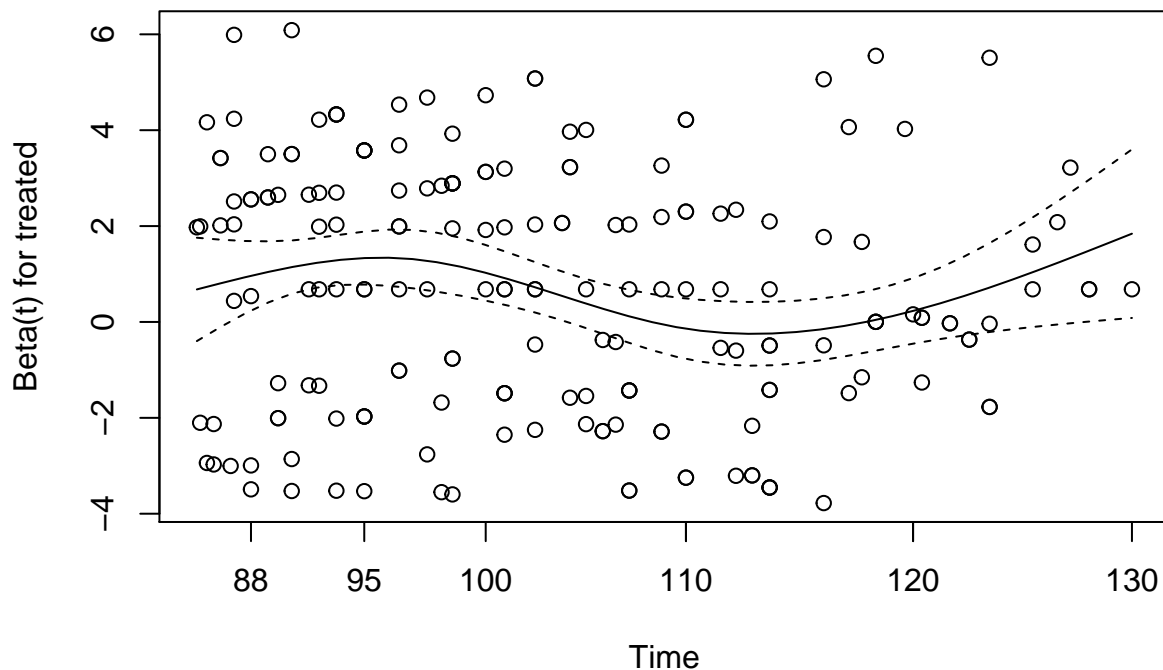
```
cox.zph(adj.s.out3) ## checking the proportional hazards assumption
```

```

      rho chisq      p
treated -0.12  3.16 0.0756

plot(cox.zph(adj.s.out3), var="treated")

```



Results So Far (After Matching and Subclassification)

These subclassification results describe the average treatment effect, while the previous analyses we have completed describe the average treatment effect on the treated. This is one reason for the meaningful

difference between the estimates. Another reason is that the balance on observed covariates is much worse after stratification in some quintiles, especially Quintile 1.

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Relative Hazard Rate)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)
After 1:1 PS Match (Match: Automated)	9.81 (6.65, 12.96)	0.143 (0.021, 0.264)	N/A	N/A
After 1:1 PS Match ("Regression" Models)	9.72 (6.55, 12.89)	N/A	1.66 (1.04, 2.62)	1.79 (1.18, 2.73)
After PS Subclassification ("Regression" models, ATE)	5.79 (2.32, 9.26)	N/A	1.94 (1.11, 9.26)	1.98 (1.41, 2.77)

Task 8. Execute weighting by the inverse PS, then assess covariate balance


ATT approach: Weight treated subjects as 1; control subjects as $ps/(1-ps)$

```
toy$wts1 <- ifelse(toy$treated==1, 1, toy$ps/(1-toy$ps))
```

Here is a plot of the resulting ATT (average treatment effect on the treated) weights:

```
ggplot(toy, aes(x = ps, y = wts1, color = treated_f)) +
  geom_point() +
  guides(color = FALSE) +
  facet_wrap(~ treated_f) +
  labs(x = "Estimated Propensity for Treatment",
       y = "ATT weights for the toy example",
       title = "ATT weighting structure: Toy example")
```

ATT weighting structure: Toy example




ATE Approach: Weight treated subjects by $1/ps$; Control subjects by $1/(1-PS)$

```
toy$wts2 <- ifelse(toy$treated==1, 1/toy$ps, 1/(1-toy$ps))
```

Here's a plot of the ATE (average treatment effect) weights...

```
ggplot(toy, aes(x = ps, y = wts2, color = treated_f)) +
  geom_point() +
  guides(color = FALSE) +
  facet_wrap(~ treated_f) +
  labs(x = "Estimated Propensity for Treatment",
       y = "ATE weights for the toy example",
       title = "ATE weighting structure: Toy example")
```

ATE weighting structure: Toy example



Assessing Balance after Weighting

The `twang` package provides several functions for assessing balance after weighting, in addition to actually doing the weighting using more complex propensity models. For this example, we'll demonstrate balance assessment for our two (relatively simple) weighting schemes. In other examples, we'll use `twang` to do more complete weighting work.

Reminder of ATT vs. ATE Definitions

- Informally, the **average treatment effect on the treated** (ATT) estimate describes the difference in potential outcomes (between treated and untreated subjects) summarized across the population of people who actually received the treatment. This is usually the estimate we work with in making causal estimates from observational studies.
- On the other hand, the **average treatment effect** (ATE) refers to the difference in potential outcomes summarized across the entire population, including those who did not receive the treatment.

For ATT weights (`wts1`)

```
toy_df <- data.frame(toy) # twang doesn't react well to tibbles

covlist <- c("covA", "covB", "covC", "covD", "covE", "covF", "Asqr", "BC", "BD", "ps", "linps")

# for ATT weights
```

```
bal.wts1 <- dx.wts(x=toy_df$wts1, data=toy_df, vars=covlist,
                  treat.var="treated", estimand="ATT")
```

```
bal.wts1
```

```

  type n.treat n.ctrl ess.treat ess.ctrl    max.es    mean.es    max.ks
1  unw      140    260      140 260.0000 1.1070181 0.43969555 0.3934066
2           140    260      140 117.3756 0.1197246 0.05601621 0.1295878

  mean.ks iter
1 0.20380389  NA
2 0.06990453  NA
```

```
bal.table(bal.wts1)
```

```
$unw
```

	tx.mn	tx.sd	ct.mn	ct.sd	std.eff.sz	stat	p	ks
covA	3.165	1.138	3.005	1.094	0.141	1.361	0.174	0.119
covB	0.514	0.502	0.296	0.457	0.435	4.284	0.000	0.218
covC	9.624	1.873	10.596	2.045	-0.519	-4.800	0.000	0.241
covD	9.159	2.083	8.647	2.212	0.246	2.300	0.022	0.116
covE	9.771	2.839	11.300	3.423	-0.538	-4.779	0.000	0.225
covF:1-Low	0.271	0.445	0.454	0.498	-0.410	9.831	0.000	0.182
covF:2-Middle	0.386	0.487	0.377	0.485	0.018	NA	NA	0.009
covF:3-High	0.343	0.475	0.169	0.375	0.366	NA	NA	0.174
Asqr	11.301	6.743	10.219	6.014	0.161	1.592	0.112	0.119
BC	4.952	5.016	2.992	4.781	0.391	3.796	0.000	0.237
BD	4.520	4.661	2.440	3.927	0.446	4.499	0.000	0.223
ps	0.459	0.179	0.291	0.185	0.939	8.879	0.000	0.393
linps	-0.189	0.801	-1.076	1.012	1.107	9.624	0.000	0.393

```
ks.pval
```

covA	0.141
covB	0.000
covC	0.000
covD	0.155
covE	0.000
covF:1-Low	0.000
covF:2-Middle	0.000
covF:3-High	0.000
Asqr	0.141
BC	0.000
BD	0.000
ps	0.000
linps	0.000

```
[[2]]
```

	tx.mn	tx.sd	ct.mn	ct.sd	std.eff.sz	stat	p	ks
covA	3.165	1.138	3.187	1.133	-0.020	-0.155	0.877	0.079
covB	0.514	0.502	0.556	0.498	-0.082	-0.675	0.500	0.041
covC	9.624	1.873	9.550	2.206	0.039	0.281	0.778	0.083
covD	9.159	2.083	9.212	1.997	-0.025	-0.212	0.833	0.062
covE	9.771	2.839	9.750	2.834	0.008	0.063	0.950	0.078
covF:1-Low	0.271	0.445	0.229	0.420	0.096	0.334	0.706	0.043
covF:2-Middle	0.386	0.487	0.409	0.492	-0.048	NA	NA	0.023
covF:3-High	0.343	0.475	0.362	0.481	-0.041	NA	NA	0.020
Asqr	11.301	6.743	11.435	6.459	-0.020	-0.157	0.876	0.079
BC	4.952	5.016	5.281	5.050	-0.066	-0.542	0.588	0.062

BD	4.520	4.661	4.860	4.574	-0.073	-0.588	0.557	0.081
ps	0.459	0.179	0.481	0.195	-0.120	-0.914	0.361	0.130
linps	-0.189	0.801	-0.116	0.904	-0.091	-0.696	0.487	0.130
	ks.pval							
covA	0.784							
covB	1.000							
covC	0.735							
covD	0.950							
covE	0.792							
covF:1-Low	0.706							
covF:2-Middle	0.706							
covF:3-High	0.706							
Asqr	0.784							
BC	0.949							
BD	0.762							
ps	0.209							
linps	0.209							

The `std.eff.sz` shows the standardized difference, but as a proportion, rather than as a percentage. We'll create a data frame (tibble) so we can plot the data more easily.

```
bal.before.wts1 <- bal.table(bal.wts1)[1]
bal.after.wts1 <- bal.table(bal.wts1)[2]


balance.att.weights <- data_frame(names = rownames(bal.before.wts1$unw),
                                   pre.weighting = 100*bal.before.wts1$unw$std.eff.sz,
                                   ATT.weighted = 100*bal.after.wts1[[1]]$std.eff.sz)
balance.att.weights <- gather(balance.att.weights, timing, szd, 2:3)
```

OK - here is the plot of standardized differences before and after ATT weighting.

```
ggplot(balance.att.weights, aes(x = szd, y = reorder(names, szd), color = timing)) +
  geom_point() +
  geom_vline(xintercept = 0) +
  geom_vline(xintercept = c(-10,10), linetype = "dashed", col = "blue") +
  labs(x = "Standardized Difference", y = "",
       title = "Standardized Difference before and after ATT Weighting",
       subtitle = "The toy example")
```

Standardized Difference before and after ATT Weighting

The toy example



For ATE weights (wts2)

```
bal.wts2 <- dx.wts(x=toy_df$wts2, data=toy_df, vars=covlist,
  treat.var="treated", estimand="ATE")
```

```
bal.wts2
```

	type	n.treat	n.ctrl	ess.treat	ess.ctrl	max.es	mean.es	max.ks
1	unw	140	260	140.0000	260.000	0.8585784	0.40894619	0.3934066
2		140	260	111.5654	224.749	0.1771696	0.07730818	0.1876510

	mean.ks	iter
1	0.20380389	NA
2	0.08075449	NA

```
bal.table(bal.wts2)
```

```
$unw
```

	tx.mn	tx.sd	ct.mn	ct.sd	std.eff.sz	stat	p	ks
covA	3.165	1.138	3.005	1.094	0.144	1.361	0.174	0.119
covB	0.514	0.502	0.296	0.457	0.451	4.284	0.000	0.218
covC	9.624	1.873	10.596	2.045	-0.477	-4.800	0.000	0.241
covD	9.159	2.083	8.647	2.212	0.235	2.300	0.022	0.116
covE	9.771	2.839	11.300	3.423	-0.462	-4.779	0.000	0.225
covF:1-Low	0.271	0.445	0.454	0.498	-0.410	9.831	0.000	0.182
covF:2-Middle	0.386	0.487	0.377	0.485	0.018	NA	NA	0.009
covF:3-High	0.343	0.475	0.169	0.375	0.366	NA	NA	0.174
Asqr	11.301	6.743	10.219	6.014	0.172	1.592	0.112	0.119

BC	4.952	5.016	2.992	4.781	0.396	3.796	0.000	0.237
BD	4.520	4.661	2.440	3.927	0.483	4.499	0.000	0.223
ps	0.459	0.179	0.291	0.185	0.844	8.879	0.000	0.393
linps	-0.189	0.801	-1.076	1.012	0.859	9.624	0.000	0.393
	ks.pval							
covA	0.141							
covB	0.000							
covC	0.000							
covD	0.155							
covE	0.000							
covF:1-Low	0.000							
covF:2-Middle	0.000							
covF:3-High	0.000							
Asqr	0.141							
BC	0.000							
BD	0.000							
ps	0.000							
linps	0.000							

[[2]]

	tx.mn	tx.sd	ct.mn	ct.sd	std.eff.sz	stat	p	ks
covA	3.146	1.105	3.070	1.111	0.069	0.602	0.548	0.082
covB	0.415	0.495	0.389	0.489	0.052	0.456	0.649	0.026
covC	10.033	1.894	10.220	2.164	-0.092	-0.791	0.429	0.113
covD	9.125	2.261	8.850	2.154	0.124	1.027	0.305	0.118
covE	10.442	2.949	10.743	3.309	-0.096	-0.847	0.398	0.084
covF:1-Low	0.345	0.475	0.373	0.484	-0.059	0.146	0.864	0.028
covF:2-Middle	0.396	0.489	0.388	0.487	0.016	NA	NA	0.008
covF:3-High	0.259	0.438	0.239	0.426	0.046	NA	NA	0.020
Asqr	11.111	6.583	10.656	6.205	0.071	0.610	0.543	0.082
BC	4.076	5.009	3.814	5.002	0.052	0.459	0.646	0.068
BD	3.550	4.470	3.310	4.330	0.055	0.478	0.633	0.045
ps	0.378	0.176	0.359	0.209	0.096	0.823	0.411	0.188
linps	-0.561	0.806	-0.731	1.078	0.177	1.569	0.118	0.188
	ks.pval							
covA	0.666							
covB	1.000							
covC	0.273							
covD	0.226							
covE	0.627							
covF:1-Low	0.864							
covF:2-Middle	0.864							
covF:3-High	0.864							
Asqr	0.666							
BC	0.849							
BD	0.996							
ps	0.009							
linps	0.009							

```
bal.before.wts2 <- bal.table(bal.wts2)[1]
```

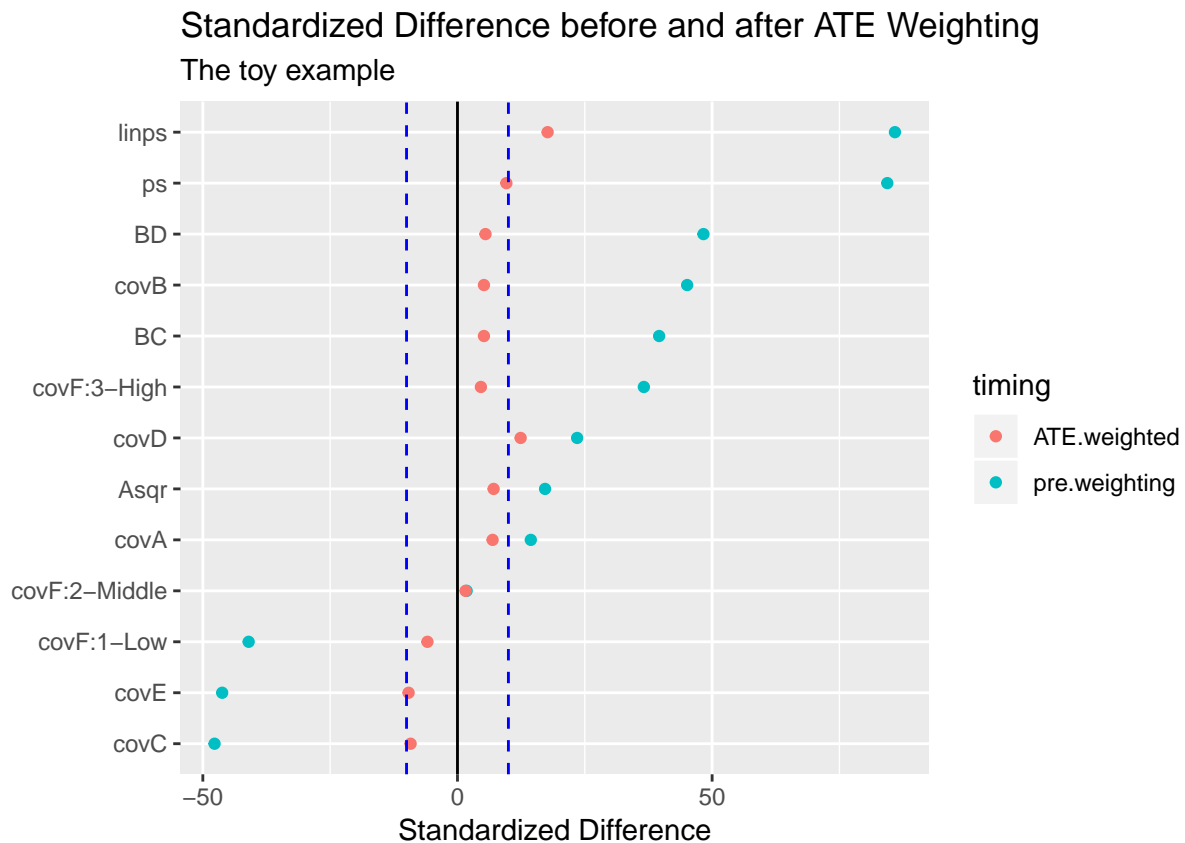
```
bal.after.wts2 <- bal.table(bal.wts2)[2]
```

```
balance.ate.weights <- data_frame(names = rownames(bal.before.wts2$unw),
                                   pre.weighting = 100*bal.before.wts2$unw$std.eff.sz,
```

```
ATE.weighted = 100*bal.after.wts2[[1]]$std.eff.sz)
balance.ate.weights <- gather(balance.ate.weights, timing, szd, 2:3)
```

Here is the plot of standardized differences before and after ATE weighting.

```
ggplot(balance.ate.weights, aes(x = szd, y = reorder(names, szd), color = timing)) +
  geom_point() +
  geom_vline(xintercept = 0) +
  geom_vline(xintercept = c(-10,10), linetype = "dashed", col = "blue") +
  labs(x = "Standardized Difference", y = "",
       title = "Standardized Difference before and after ATE Weighting",
       subtitle = "The toy example")
```



Rubin's Rules after ATT weighting

For our weighted sample, our summary statistic for Rules 1 and 2 may be found from the `bal.table` output.

Rubin's Rule 1

We can read off the standardized effect size after weighting for the linear propensity score as -0.091. Multiplying by 100, we get 9.1%, so we would pass Rule 1.

Rubin's Rule 2

We can read off the standard deviations within the treated and control groups. We can then square each, to get the relevant variances, then take the ratio of those variances. Here, we have standard deviations of the linear propensity score after weighting of 0.801 in the treated group and 0.904 in the control group. $0.801^2 / 0.904^2 = 0.7851$, which is just outside our desired range of $4/5$ to $5/4$, as well as clearly within $1/2$ to 2 . Arguably, we can pass Rule 2, also. But I'll be interested to see if **twang** can do better.

Rubin's Rule 3

Rubin's Rule 3 requires some more substantial manipulation of the data. I'll skip that for now.

Rubin's Rules after ATE weighting

Again, our summary statistic for Rules 1 and 2 may be found from the `bal.table` output.

Rubin's Rule 1

The standardized effect size after ATE weighting for the linear propensity score is 0.177. Multiplying by 100, we get 17.7%, so we would pass Rule 1.

Rubin's Rule 2

We can read off the standard deviations within the treated and control groups from the ATE weights, then square to get the variances, then take the ratio. Here, we have $0.806^2 / 1.078^2 = 0.559$, which is not within our desired range of $4/5$ to $5/4$, but is between 0.5 and 2. Arguably, we pass Rule 2, also. But I'll be interested to see if **twang** can do better.

Rubin's Rule 3

Again, for now, I'm skipping Rubin's Rule 3 after weighting.

Using TWANG for Alternative PS Estimation and ATT Weighting

Here, I'll demonstrate the use of the **twang** package's functions to fit the propensity model and then perform ATT weighting, mostly using default options.

Estimate the Propensity Score using Generalized Boosted Regression, and then perform ATT Weighting

We can directly use the **twang** (toolkit for **w**eighting and **a**nalysis of **n**onequivalent **g**roups) package to weight our results, and even to re-estimate the propensity score using generalized boosted regression rather than a logistic regression model. The **twang** vignette is very helpful and found at this link.


To begin, we'll estimate the propensity score using the **twang** function `ps`. This uses a *generalized boosted regression* approach to estimate the propensity score and produce material for checking balance.

```
# Recall that twang does not play well with tibbles,
# so we have to use the data frame version of the toy object

ps.toy <- ps(treated ~ covA + covB + covC + covD + covE + covF +
             Asqr + BC + BD,
             data = toy_df,
             n.trees = 3000,
             interaction.depth = 2,
             stop.method = c("es.mean"),
             estimand = "ATT",
             verbose = FALSE)
```

Did we let the simulations run long enough to stabilize estimates?

```
plot(ps.toy)
```



What is the effective sample size of our weighted results?

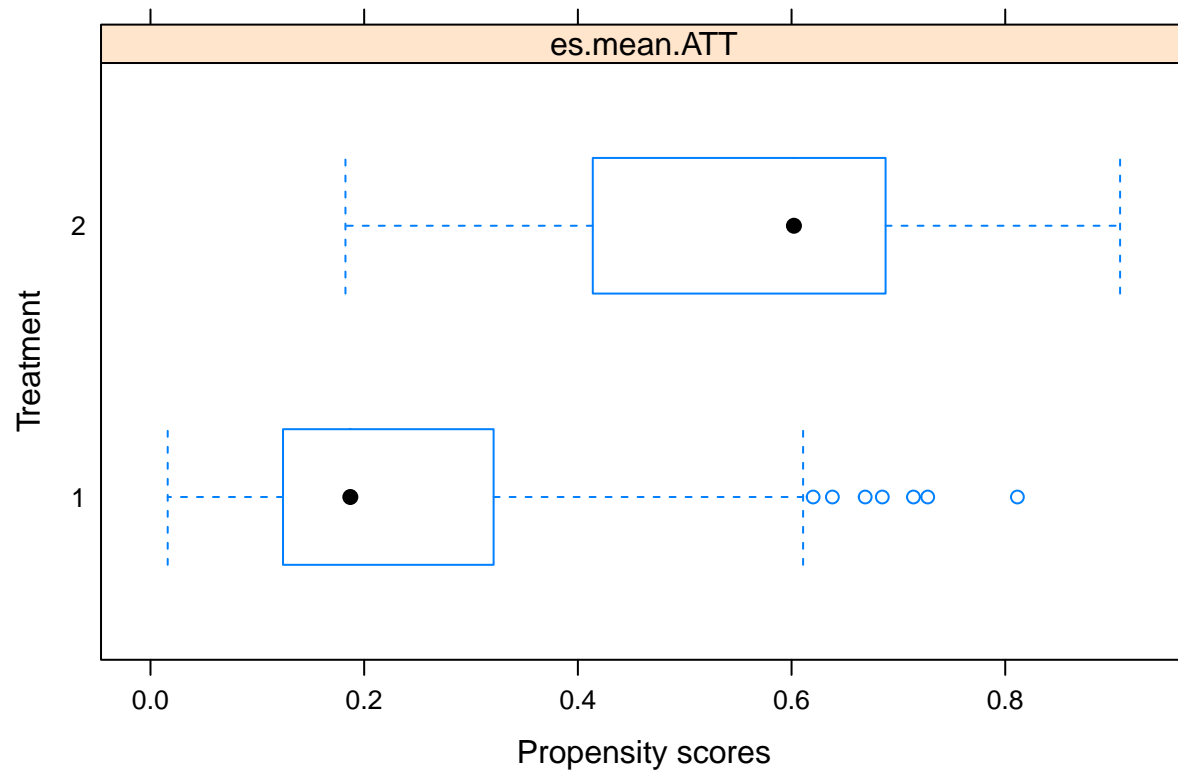
```
summary(ps.toy)
```

	n.treat	n.ctrl	ess.treat	ess.ctrl	max.es	mean.es
unw	140	260	140	260.0000	0.53833327	0.33365329
es.mean.ATT	140	260	140	107.1175	0.08192421	0.04338572

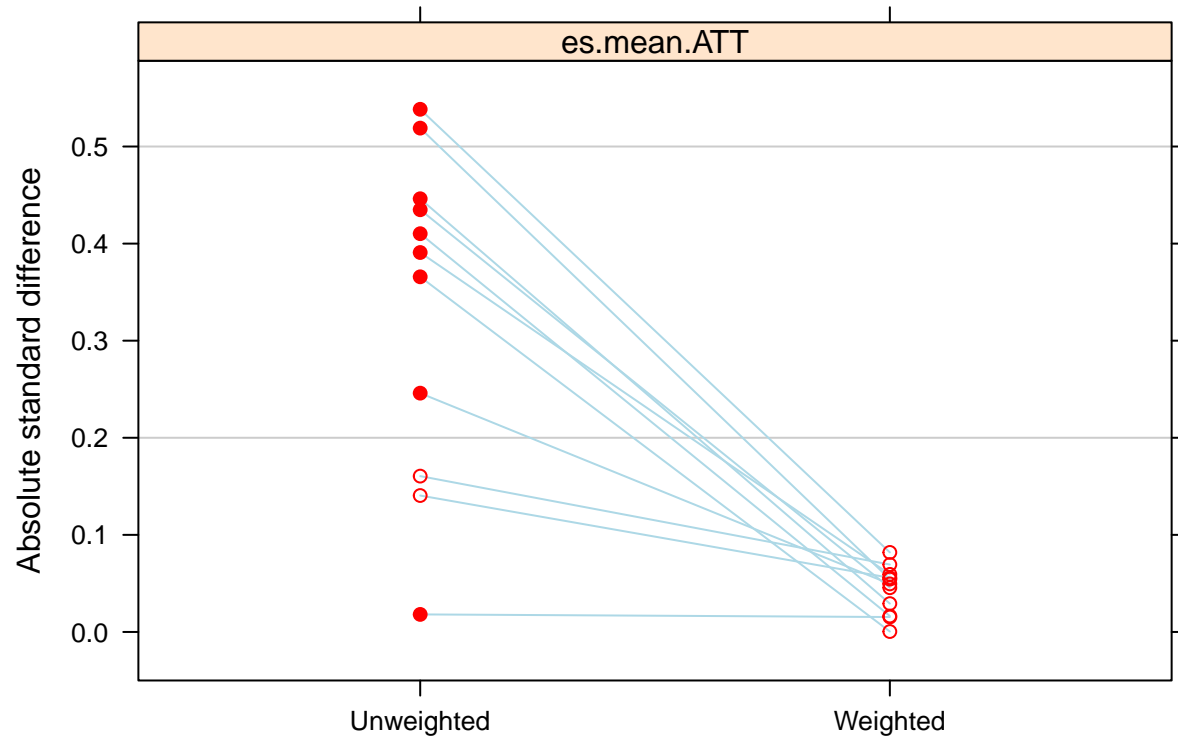
	max.ks	max.ks.p	mean.ks	iter
unw	0.24065934	NA	0.16933067	NA
es.mean.ATT	0.07831191	NA	0.04627681	1128

How is the balance?

```
plot(ps.toy, plots = 2)
```



```
plot(ps.toy, plots = 3)
```



Assessing Balance with cobalt

```
bal.tab(ps.toy, full.stop.method = "es.mean.att")
```

Call

```
ps(formula = treated ~ covA + covB + covC + covD + covE + covF +
  Asqr + BC + BD, data = toy_df, n.trees = 3000, interaction.depth = 2,
  verbose = FALSE, estimand = "ATT", stop.method = c("es.mean"))
```

Balance Measures


	Type	Diff.Adj
prop.score	Distance	0.9594
covA	Contin.	0.0558
covB	Binary	0.0228
covC	Contin.	-0.0544
covD	Contin.	-0.0495
covE	Contin.	-0.0819
covF_1-Low	Binary	0.0073
covF_2-Middle	Binary	-0.0075
covF_3-High	Binary	0.0002
Asqr	Contin.	0.0694
BC	Contin.	0.0593
BD	Contin.	0.0292

Effective sample sizes

	Control	Treated
Unadjusted	260.000	140
Adjusted	107.117	140

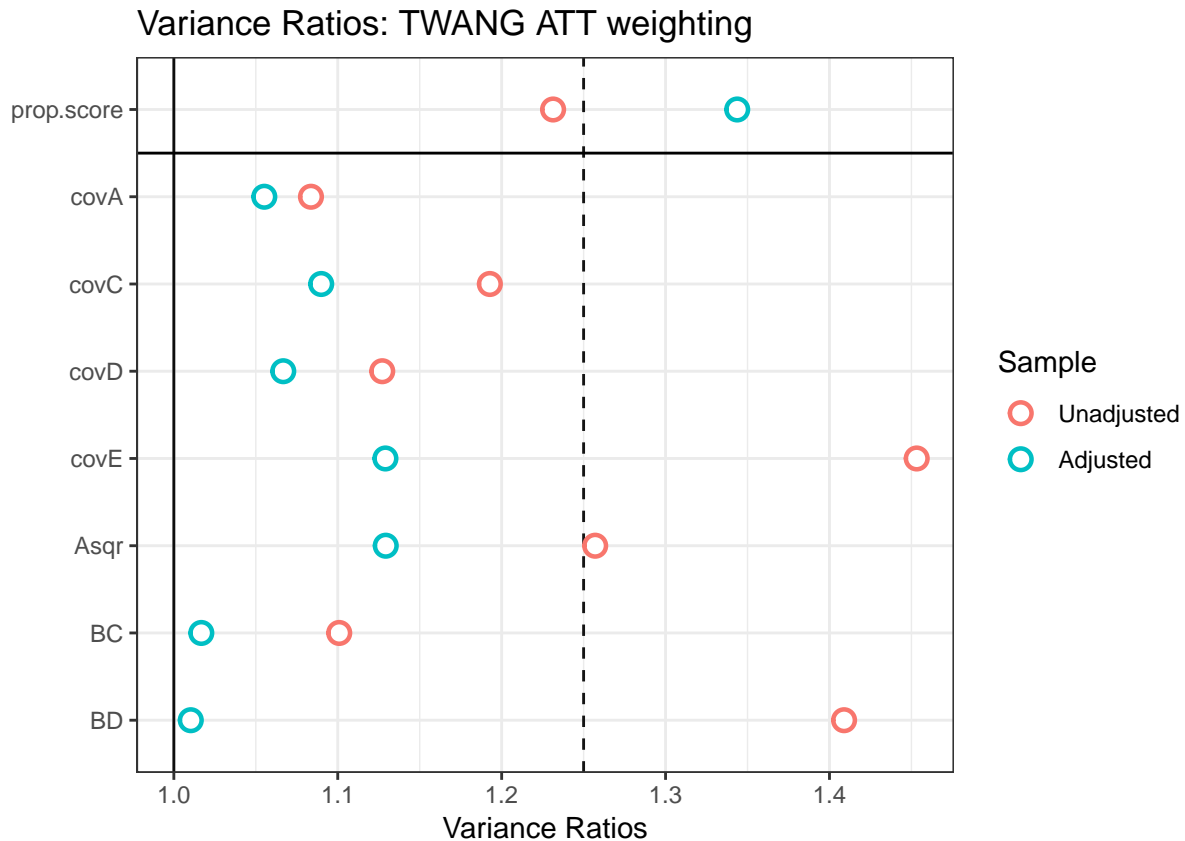
Semi-Automated Love plot of Standardized Differences

```
p <- love.plot(bal.tab(ps.toy),
               threshold = .1, size = 1.5,
               title = "Standardized Diffs and TWANG ATT weighting")
p + theme_bw()
```



Semi-Automated Love plot of Variance Ratios

```
p <- love.plot(bal.tab(ps.toy), stat = "v",
               threshold = 1.25, size = 1.5,
               title = "Variance Ratios: TWANG ATT weighting")
p + theme_bw()
```



Task 9. After weighting, what is the estimated average causal effect of treatment?

... on Outcome 1 [a continuous outcome]

with ATT weights

The relevant regression approach uses the `svydesign` and `svyglm` functions from the `survey` package.

```
toywt1.design <- svydesign(ids=~1, weights=~wts1, data=toy) # using ATT weights
adjout1.wt1 <- svyglm(out1.cost ~ treated, design=toywt1.design)
wt_att_results1 <- tidy(adjout1.wt1, conf.int = TRUE) %>% filter(term == "treated")
```

with ATE weights

```
toywt2.design <- svydesign(ids=~1, weights=~wts2, data=toy) # using ATE weights
adjout1.wt2 <- svyglm(out1.cost ~ treated, design=toywt2.design)
wt_ate_results1 <- tidy(adjout1.wt2, conf.int = TRUE) %>% filter(term == "treated")
```

with TWANG ATT weights

```
toywt3.design <- svydesign(ids=~1,
                          weights=~get.weights(ps.toy,
                                                stop.method = "es.mean"),
                          data=toy) # using twang ATT weights

adjout1.wt3 <- svyglm(out1.cost ~ treated, design=toywt3.design)
wt_twangatt_results1 <- tidy(adjout1.wt3, conf.int = TRUE) %>% filter(term == "treated")
```

... on Outcome 2 [a binary outcome]

For a binary outcome, we build the outcome model using the quasibinomial, rather than the usual binomial family. We use the same `svydesign` information as we built for outcome 1.

Using ATT weights

```
adjout2.wt1 <- svyglm(out2 ~ treated, design=toywt1.design, family=quasibinomial())

wt_att_results2 <- tidy(adjout2.wt1, conf.int = TRUE, exponentiate = TRUE) %>%
  filter(term == "treated")
```

Using ATE weights

```
adjout2.wt2 <- svyglm(out2.event ~ treated, design=toywt2.design, family=quasibinomial())

wt_ate_results2 <- tidy(adjout2.wt2, conf.int = TRUE, exponentiate = TRUE) %>%
  filter(term == "treated")
```

with TWANG ATT weights

```
adjout2.wt3 <- svyglm(out2 ~ treated, design=toywt3.design,
                      family=quasibinomial())

wt_twangatt_results2 <- tidy(adjout2.wt3, conf.int = TRUE, exponentiate = TRUE) %>%
  filter(term == "treated")
```

... on Outcome 3 [a time to event]

As before, subjects with `out2.event = "Yes"` are truly observed events, while those with `out2.event == "No"` are censored before an event can happen to them.

Using ATT weights

The Cox model comparing treated to control, weighting by ATT weights (`wts1`), is...

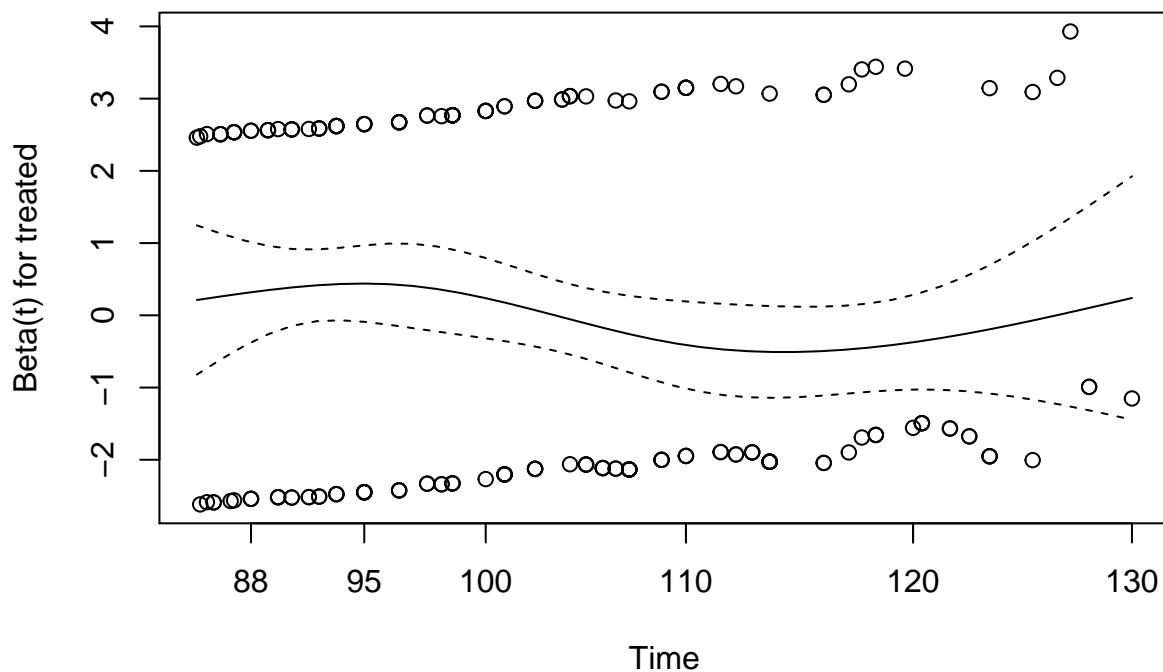
```
adjout3.wt1 <- coxph(Surv(out3.time, out2) ~ treated, data=toy, weights=wts1)
wt_att_results3 <- tidy(adjout3.wt1, exponentiate = TRUE) %>%
  filter(term == "treated")
```

The `exp(coef)` output gives the relative hazard of the event comparing treated subjects to control subjects.

And here's the check of the proportional hazards assumption...

```
cox.zph(adjout3.wt1); plot(cox.zph(adjout3.wt1), var="treated")
```

```
      rho chisq      p
treated -0.109   2.7 0.101
```



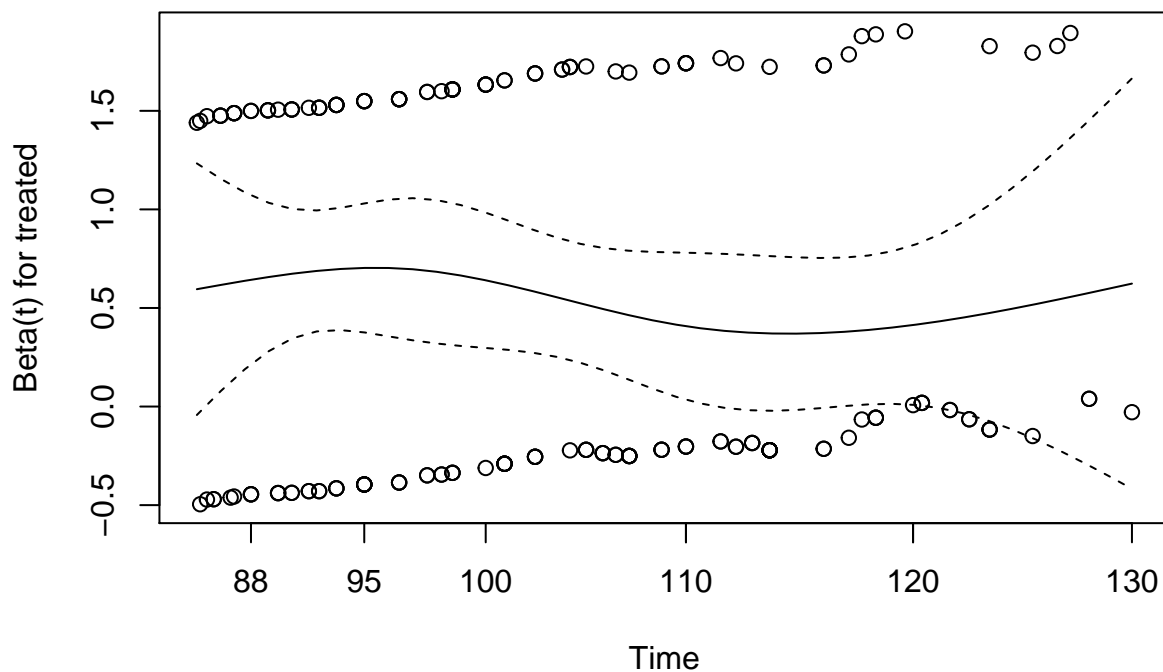
Using ATE weights

```
adjout3.wt2 <- coxph(Surv(out3.time, out2) ~ treated, data=toy, weights=wts2)
wt_ate_results3 <- tidy(adjout3.wt2, exponentiate = TRUE) %>%
  filter(term == "treated")
```

And here's the check of the proportional hazards assumption...

```
cox.zph(adjout3.wt2); plot(cox.zph(adjout3.wt2), var="treated")
```

```
      rho chisq      p
treated -0.0974 0.815 0.367
```

with TWANG ATT weights

```
wt3 <- get.weights(ps.toy, stop.method = "es.mean")

adjout3.wt3 <- coxph(Surv(out3.time, out2) ~ treated, data=toy, weights=wt3)
wt_twangatt_results3 <- tidy(adjout3.wt3, exponentiate = TRUE) %>%
  filter(term == "treated")
```

Results So Far (After Matching, Subclassification and Weighting)

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Rel. HR)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)
After 1:1 PS Match (Match: Automated)	9.81 (6.65, 12.96)	0.143 (0.021, 0.264)	N/A	N/A
After 1:1 PS Match ("Regression" Models)	9.72 (6.55, 12.89)	N/A	1.66 (1.04, 2.62)	1.79 (1.18, 2.73)
After PS Subclassification ("Regression" models, ATE)	5.79 (2.32, 9.26)	N/A	1.94 (1.11, 9.26)	1.98 (1.41, 2.77)

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Rel. HR)
ATT Weighting (ATT)	7.70 (4.02, 11.38)	N/A N/A	1.59 (0.97, 2.61)	1.76 (1.27, 2.42)
ATE Weighting (ATE)	7.44 (4.15, 10.73)	N/A N/A	2.06 (1.30, 3.27)	2.22 (1.82, 2.71)
twang ATT weights (ATT)	8.40 (5.01, 11.78)	N/A N/A	1.57 (0.95, 2.61)	1.76 (1.23, 2.51)

Task 10. After direct adjustment for the linear PS, what is the estimated average causal treatment effect?

... on Outcome 1 [a continuous outcome]

Here, we fit a linear regression model with `linps` added as a covariate.

```
adj.reg.out1 <- lm(out1.cost ~ treated + linps, data=toy)
adj_out1 <- tidy(adj.reg.out1, conf.int = TRUE) %>% filter(term == "treated")
```

... on Outcome 2 [a binary outcome]

Here, fit a logistic regression with `linps` added as a covariate

```
adj.reg.out2 <- glm(out2 ~ treated + linps, data=toy, family=binomial())
adj_out2 <- tidy(adj.reg.out2, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term == "treated")
```

... on Outcome 3 [a time-to-event outcome]

Again, subjects with `out2.event` No are right-censored, those with Yes for `out2.event` have their times to event observed.

We fit a Cox proportional hazards model predicting time to event (with `event=Yes` indicating non-censored cases) based on treatment group (`treated`) and now also the linear propensity score.

```
adj.reg.out3 <- coxph(Surv(out3.time, out2) ~ treated + linps, data=toy)
adj_out3 <- tidy(adj.reg.out3, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term == "treated")
```

The `exp(coef)` section of the `summary` for this model indicates the relative hazard estimates and associated 95% CI.

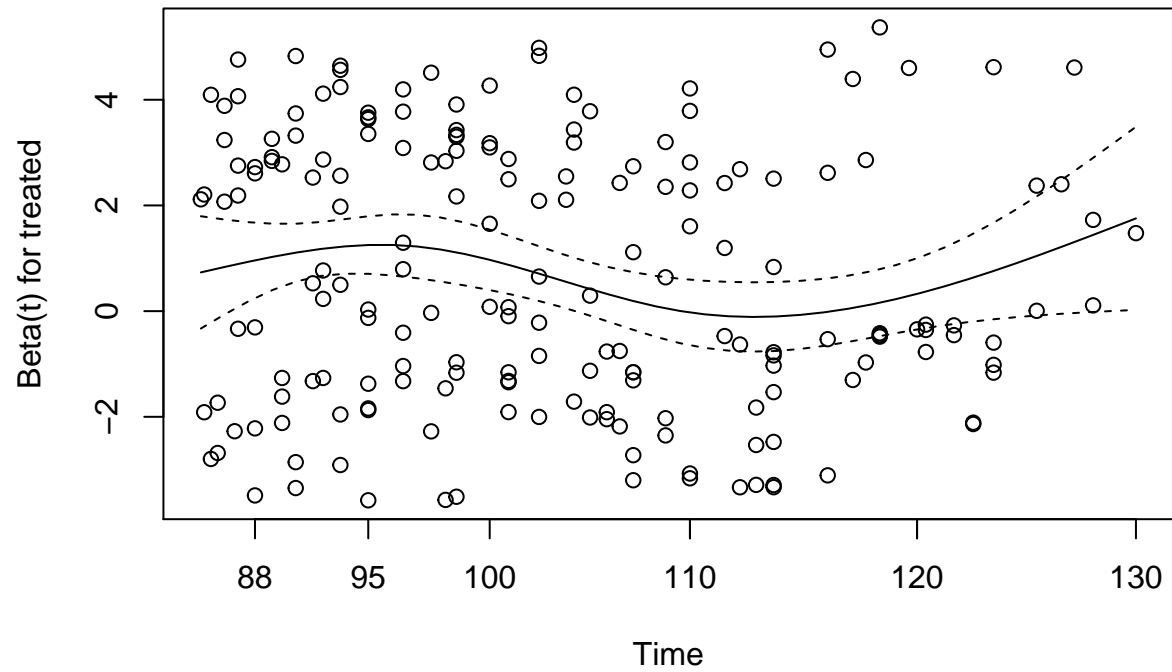
Check proportional hazards assumption

Here's the check of the proportional hazards assumption.


```
cox.zph(adj.reg.out3)
```

	rho	chisq	p
treated	-0.1042	2.37	0.124
linps	0.0894	1.57	0.210
GLOBAL	NA	2.72	0.256

```
plot(cox.zph(adj.reg.out3), var="treated")
```



```
plot(cox.zph(adj.reg.out3), var="linps")
```



Results So Far (After Matching, Subclassification, Weighting, Adjustment)

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Rel. HR)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)
After 1:1 PS Match (Match: Automated)	9.81 (6.65, 12.96)	0.143 (0.021, 0.264)	N/A	N/A
After 1:1 PS Match ("Regression" Models)	9.72 (6.55, 12.89)	N/A	1.66 (1.04, 2.62)	1.79 (1.18, 2.73)
After PS Subclassification ("Regression" models, ATE)	5.79 (2.32, 9.26)	N/A	1.94 (1.11, 9.26)	1.98 (1.41, 2.77)
ATT Weighting (ATT)	7.70 (4.02, 11.38)	N/A	1.59 (0.97, 2.61)	1.76 (1.27, 2.42)
ATE Weighting (ATE)	7.44 (4.15, 10.73)	N/A	2.06 (1.30, 3.27)	2.22 (1.82, 2.71)
twang ATT weights (ATT)	8.40 (5.01, 11.78)	N/A	1.57 (0.95, 2.61)	1.76 (1.23, 2.51)
Direct Adjustment (with linps, ATT)	7.99 (4.86, 11.13)	N/A	1.80 (1.14, 2.85)	1.80 (1.14, 2.85)

Task 11. “Double Robust” Approach - Weighting + Adjustment, what is the estimated average causal effect of treatment?

This approach is essentially identical to the weighting analyses done in Task 9. The only change is to add `linps` to `treated` in the outcome models.

... on Outcome 1 [a continuous outcome]

with ATT weights

The relevant regression approach uses the `svydesign` and `svyglm` functions from the `survey` package.

```
toywt1.design <- svydesign(ids=~1, weights=~wts1, data=toy) # using ATT weights

dr.out1.wt1 <- svyglm(out1.cost ~ treated + linps, design=toywt1.design)

dr_att_out1 <- tidy(dr.out1.wt1, conf.int = TRUE) %>% filter(term == "treated")
dr_att_out1
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic    p.value conf.low conf.high
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 treated      7.91      1.84      4.29 0.0000221     4.30     11.5
```

with ATE weights

```
toywt2.design <- svydesign(ids=~1, weights=~wts2, data=toy) # using ATE weights

dr.out1.wt2 <- svyglm(out1.cost ~ treated + linps, design=toywt2.design)

dr_ate_out1 <- tidy(dr.out1.wt2, conf.int = TRUE) %>% filter(term == "treated")
dr_ate_out1
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic    p.value conf.low conf.high
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 treated      7.01      1.69      4.15 0.0000414     3.70     10.3
```

with twang based ATT weights

```
wts3 <- get.weights(ps.toy, stop.method = "es.mean")

toywt3.design <- svydesign(ids=~1, weights=~wts3, data=toy) # twang ATT weights

dr.out1.wt3 <- svyglm(out1.cost ~ treated + linps, design=toywt3.design)

dr_twangatt_out1 <- tidy(dr.out1.wt3, conf.int = TRUE) %>% filter(term == "treated")
dr_twangatt_out1
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic    p.value conf.low conf.high
```

	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1 treated		8.14	1.72	4.72	0.00000330	4.76	11.5

... on Outcome 2 [a binary outcome]

For a binary outcome, we build the outcome model using the quasibinomial, rather than the usual binomial family. We use the same `svydesign` information as we built for outcome 1.

Using ATT weights

```
dr.out2.wt1 <- svyglm(out2 ~ treated + linps, design=toywt1.design,
                     family=quasibinomial())
dr_att_out2 <- tidy(dr.out2.wt1, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term == "treated")
dr_att_out2
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 treated    1.59     0.249     1.86  0.0639    0.975    2.59
```

Using ATE weights

```
dr.out2.wt2 <- svyglm(out2.event ~ treated + linps, design=toywt2.design,
                     family=quasibinomial())
dr_ate_out2 <- tidy(dr.out2.wt2, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term == "treated")
dr_ate_out2
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 treated    2.03     0.243     2.91  0.00379    1.26    3.27
```

Using twang ATT weights

```
dr.out2.wt3 <- svyglm(out2 ~ treated + linps, design=toywt3.design,
                     family=quasibinomial())
dr_twangatt_out2 <- tidy(dr.out2.wt3, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term == "treated")
dr_twangatt_out2
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 treated    1.58     0.265     1.74  0.0828    0.943    2.66
```

... on Outcome 3 [a time to event]

As before, subjects with `out2.event = "Yes"` are truly observed events, while those with `out2.event == "No"` are censored before an event can happen to them.

Using ATT weights

The Cox model comparing treated to control, weighting by ATT weights (`wts1`), is...

```
dr.out3.wt1 <- coxph(Surv(out3.time, out2) ~ treated + linps, data=toy, weights=wts1)
dr_att_out3 <- tidy(dr.out3.wt1, exponentiate = TRUE) %>%
  filter(term == "treated")
dr_att_out3
```

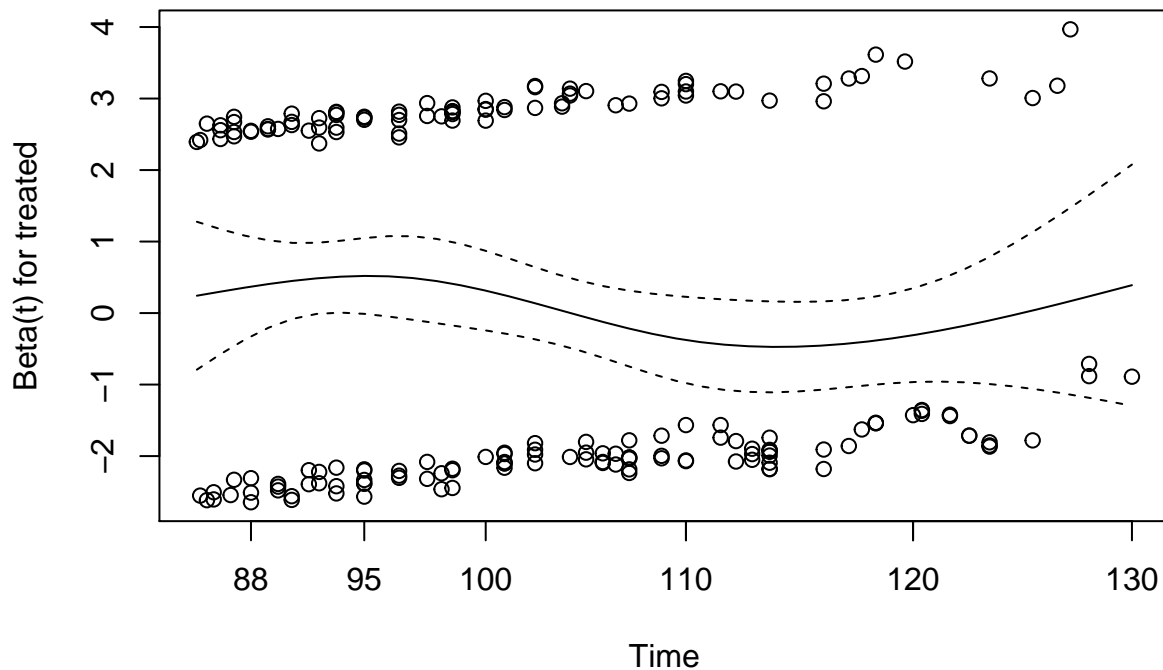
```
# A tibble: 1 x 7
  term      estimate std.error statistic  p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated      1.76      0.165      3.43 0.000613    1.27    2.43
```

The `exp(coef)` output gives the relative hazard of the event comparing treated subjects to control subjects.

And here's the check of the proportional hazards assumption...

```
cox.zph(dr.out3.wt1); plot(cox.zph(dr.out3.wt1), var="treated")
```

	rho	chisq	p
treated	-0.10962	2.63859	0.104
linps	-0.00573	0.00985	0.921
GLOBAL	NA	2.66721	0.264



Using ATE weights

```
dr.out3.wt2 <- coxph(Surv(out3.time, out2) ~ treated + linps, data=toy, weights=wts2)

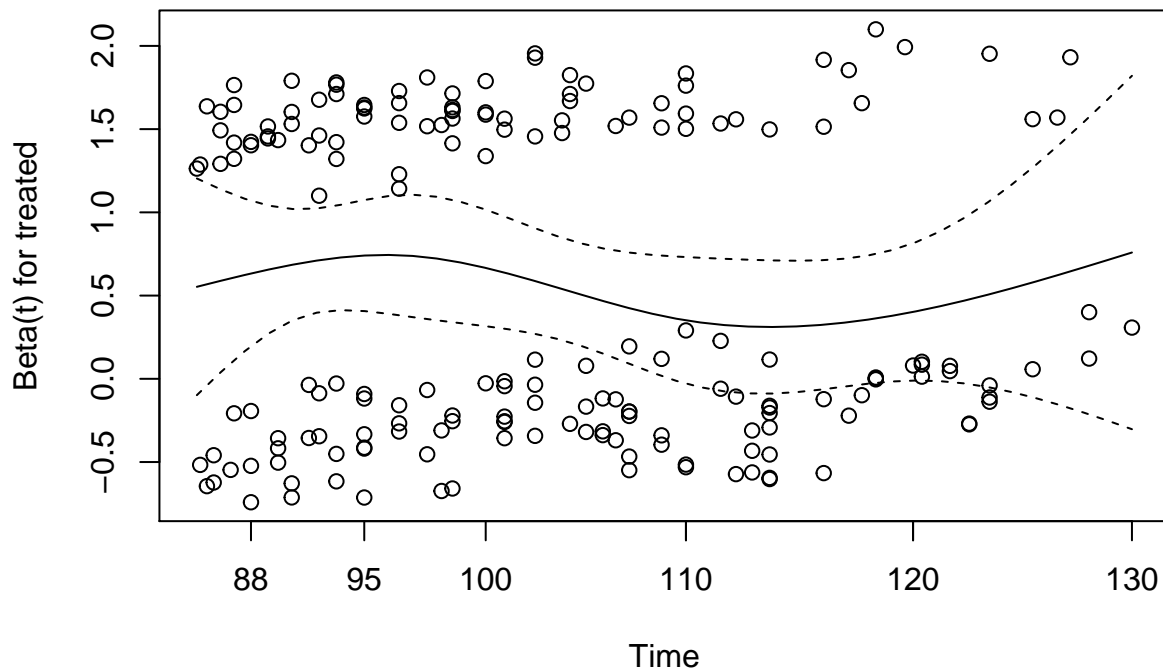
dr_ate_out3 <- tidy(dr.out3.wt2, exponentiate = TRUE) %>%
  filter(term == "treated")
dr_ate_out3
```

```
# A tibble: 1 x 7
  term      estimate std.error statistic  p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated      2.22      0.104      7.70 1.35e-14    1.81    2.72
```

And here's the check of the proportional hazards assumption...

```
cox.zph(dr.out3.wt2); plot(cox.zph(dr.out3.wt2), var="treated")
```

```
      rho chisq    p
treated -0.1030 0.865 0.352
linps    0.0249 0.055 0.815
GLOBAL    NA 0.867 0.648
```

Using twang ATT weights

```
dr.out3.wt3 <- coxph(Surv(out3.time, out2) ~ treated + linps,
                     data=toy, weights=wts3)
dr_twangatt_out3 <- tidy(dr.out3.wt3, exponentiate = TRUE) %>%
  filter(term == "treated")
dr_twangatt_out3
```

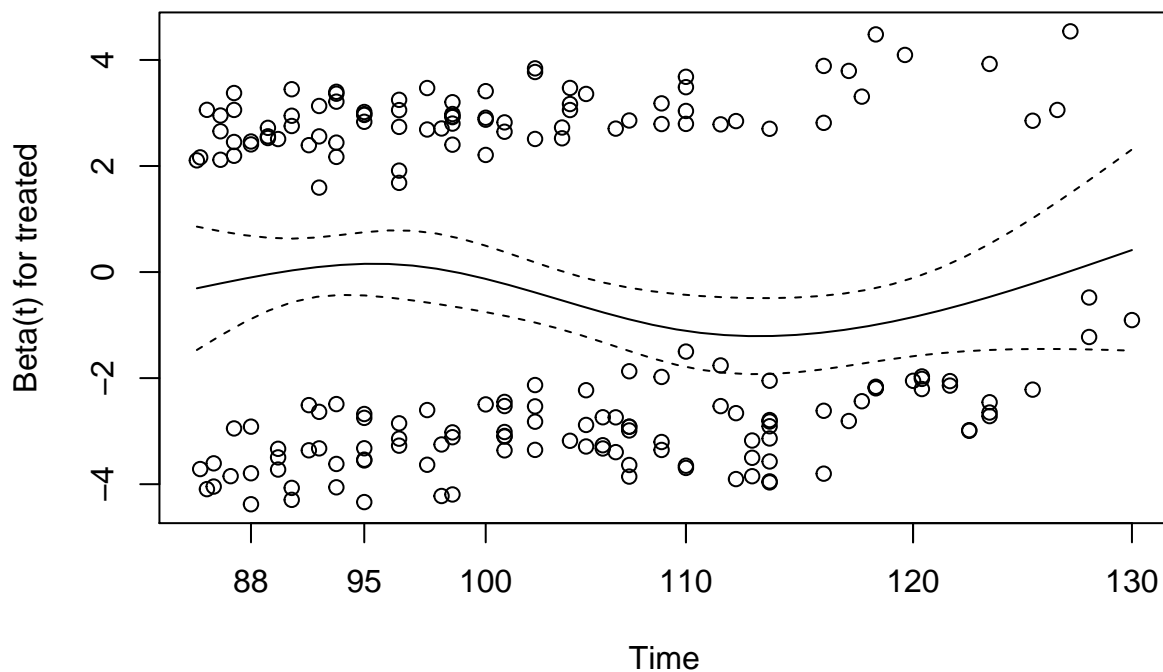
```
# A tibble: 1 x 7
  term      estimate std.error statistic p.value conf.low conf.high
<chr>      <dbl>     <dbl>     <dbl>   <dbl>   <dbl>   <dbl>
1 treated    1.80      0.185      3.17 0.00154    1.25    2.59
```

The `exp(coef)` output gives the relative hazard of the event comparing treated subjects to control subjects.

And here's the check of the proportional hazards assumption...

```
cox.zph(dr.out3.wt3); plot(cox.zph(dr.out3.wt3), var="treated")
```

```
      rho chisq    p
treated -0.1002 2.647 0.104
linps    0.0229 0.165 0.685
GLOBAL      NA 2.670 0.263
```



Task 12. Results

Treatment Effect Estimates

We now can build the table of all of the outcome results we've obtained here.

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Rel. HR)
No covariate adjustment (unadjusted)	9.64 (6.75, 12.52)	0.178 (0.075, 0.275)	2.05 (1.36, 3.13)	2.17 (1.62, 2.90)
After 1:1 PS Match (Match: Automated)	9.81 (6.65, 12.96)	0.143 (0.021, 0.264)	N/A	N/A
After 1:1 PS Match ("Regression" Models)	9.72 (6.55, 12.89)	N/A	1.66 (1.04, 2.62)	1.79 (1.18, 2.73)
After PS Subclassification ("Regression" models, ATE)	5.79 (2.32, 9.26)	N/A	1.94 (1.11, 9.26)	1.98 (1.41, 2.77)
ATT Weighting (ATT)	7.70 (4.02, 11.38)	N/A	1.59 (0.97, 2.61)	1.76 (1.27, 2.42)
ATE Weighting (ATE)	7.44 (4.15, 10.73)	N/A	2.06 (1.30, 3.27)	2.22 (1.82, 2.71)
twang ATT weights	8.40	N/A	1.57	1.76

Est. Treatment Effect (95% CI)	Outcome 1 (Cost diff.)	Outcome 2 (Risk diff.)	Outcome 2 (Odds Ratio)	Outcome 3 (Rel. HR)
(ATT)	(5.01, 11.78)	N/A	(0.95, 2.61)	(1.23, 2.51)
Direct Adjustment	7.99	N/A	1.80	1.80
(with <code>linps</code> , ATT)	(4.86, 11.13)	N/A	(1.14, 2.85)	(1.14, 2.85)
Double Robust	7.91	N/A	1.59	1.76
(ATT wts + adj.)	(4.30, 11.53)	N/A	(0.97, 2.59)	(1.27, 2.43)
Double Robust	7.01	N/A	2.03	2.22
(ATE wts + adj.)	(3.70, 10.32)	N/A	(1.26, 3.27)	(1.81, 2.72)
Double Robust	8.14	N/A	1.58	1.80
(<code>twang</code> ATT wts + adj.)	(4.76, 11.52)	N/A	(0.94, 2.66)	(1.25, 2.59)

So, with the exception of the subclassification approach (which was problematic in terms of observed covariate balance) we observe significant results (indicating higher costs with the treatment, and higher likelihood of experiencing the event, and increased hazard of event occurrence) for every adjustment approach.

Quality of Balance: Standardized Differences and Variance Ratios

We're looking at the balance across the following 10 covariates and transformations here: `covA`, `covB`, `covC`, `covD`, `covE`, `covF[middle]`, `covF[high]`, `A_squared`, `BxC` and `BxD`, as well as the raw and linear propensity scores ...

Approach	Standardized Diffs	Variance Ratios
Most Desirable Values	Between -10 and +10	Between 0.8 and 1.25
No Adjustments	-50 to 97	0.63 to 1.61
1:1 Propensity Matching	-12 to 25	0.87 to 1.24
Subclassification Quintile 1	-161 to 212	not calculated above
Quintile 2	-32 to 71	not calculated above
Quintile 3	-33 to 60	not calculated above
Quintile 4	-32 to 10	not calculated above
Quintile 5	-32 to 17	not calculated above
Propensity Weighting, ATT	-12 to 10	0.72 to 1.12
Propensity Weighting, ATE	-10 to 18	0.56 to 1.13

Quality of Balance: Rubin's Rules

Approach	Rubin 1	Rubin 2	Rubin 3
"Pass" Range, per Rubin	0 to 50	0.5 to 2.0	0.5 to 2.0
No Adjustments	85.9	0.63	0.72 to 1.76
1:1 Propensity Matching	25.4	1.24	0.79 to 1.28
Subclassification: Quintile 1	125.8	0.01	0.00 to 0.96
Quintile 2	14.8	2.17	0.42 to 11.33
Quintile 3	22.1	1.05	0.41 to 2.20
Quintile 4	4.4	0.93	0.56 to 1.23
Quintile 5	0.2	1.60	0.56 to 1.46
Propensity Weighting, ATT	-9.1	0.79	Not evaluated
Propensity Weighting, ATE	17.7	0.56	Not evaluated

Clearly, the matching and propensity weighting show improvement over the initial (no adjustments) results, although neither is completely satisfactory in terms of all covariates. In practice, I would be comfortable with either a 1:1 match or a weighting approach, I think. It isn't likely that the subclassification will get us anywhere useful in terms of balance. Rubin's Rule 3 could also be applied after weighting on the propensity score.

What is a Sensitivity Analysis for Matched Samples?

We'll study a formal sensitivity analysis approach for **matched** samples. Note well that this specific approach is appropriate only when we have

1. a statistically significant conclusion
2. from a matched samples analysis using the propensity score.

Goal of a Formal Sensitivity Analysis for Matched Samples

To replace a general qualitative statement that applies in all observational studies, like ...

the association we observe between treatment and outcome does not imply causation

or

hidden biases can explain observed associations

... with a quantitative statement that is specific to what is observed in a particular study, such as ...

to explain the association seen in a particular study, one would need a hidden bias of a particular magnitude.

If the association is strong, the hidden bias needed to explain it would be large.

- If a study is free of hidden bias (main example: a carefully randomized trial), this means that any two units (patients, subjects, whatever) that appear similar in terms of their observed covariates actually have the same chance of assignment to treatment.
- There is *hidden bias* if two units with the same observed covariates have different chances of receiving the treatment.

A **sensitivity analysis** asks: How would inferences about treatment effects be altered by hidden biases of various magnitudes? How large would these differences have to be to alter the qualitative conclusions of the study?

The methods for building such sensitivity analyses are largely due to Paul Rosenbaum, and as a result the methods are sometimes referred to as **Rosenbaum bounds**.

The Sensitivity Parameter, Γ

Suppose we have two units (subjects, patients), say, j and k , with the same observed covariate values \mathbf{x} but different probabilities p of treatment assignment (possibly due to some unobserved covariate), so that $\mathbf{x}_j = \mathbf{x}_k$ but that possibly $p_j \neq p_k$.

Units j and k might be *matched* to form a matched pair in our attempt to control overt bias due to the covariates \mathbf{x} .

- The odds that units j and k receive the treatment are, respectively, $\frac{p_j}{1-p_j}$ and $\frac{p_k}{1-p_k}$, and the odds ratio is thus the ratio of these odds.

Imagine that we knew that this odds ratio for units with the same \mathbf{x} was at most some number Γ , so that $\Gamma \geq 1$. That is,

$$\frac{1}{\Gamma} \leq \frac{p_j(1-p_j)}{p_k(1-p_k)} \leq \Gamma$$

We call Γ the **sensitivity parameter**, and it is the basis for our sensitivity analyses.

- If $\Gamma = 1$, then $p_j = p_k$ whenever $\mathbf{x}_j = \mathbf{x}_k$, so the study would be free of hidden bias, and standard statistical methods designed for randomized trials would apply.

If $\Gamma = 2$, then two units who appear similar in that they have the same set of observed covariates \mathbf{x} , could differ in their odds of receiving the treatment by as much as a factor of 2, so that one could be twice as likely as the other to receive the treatment.

So Γ is a value between 1 and ∞ where the size of Γ indicates the degree of a departure from a study free of hidden bias.

Interpreting the Sensitivity Parameter, Γ

Again, Γ is a measure of the degree of departure from a study that is free of hidden bias.

A sensitivity analysis will consider possible values of Γ and show how the inference for our outcomes might change under different levels of hidden bias, as indexed by Γ .

- A study is *sensitive* if values of Γ close to 1 could lead to inferences that are very different from those obtained assuming the study is free of hidden bias.
- A study is *insensitive* (a good thing here) if extreme values of Γ are required to alter the inference.

When we perform this sort of sensitivity analysis, we will specify different levels of hidden bias (different Γ values) and see how large a Γ we can have while still retaining the fundamental conclusions of the matched outcomes analysis.

Task 13. Sensitivity Analysis for Matched Samples, Outcome 1, using `rbounds`

In our matched sample analysis, for outcome 1 (cost) in the toy example, we saw a statistically significant result. A formal *sensitivity analysis* is called for, as a result, and we will accomplish one for this quantitative outcome, using the `rbounds` package.

The `rbounds` package is designed to work with the output from `Matching`, and can calculate Rosenbaum sensitivity bounds for the treatment effect, which help us understand the impact of hidden bias needed to invalidate our significant conclusions from the matched samples analysis.

Rosenbaum Bounds for the Wilcoxon Signed Rank test (Quantitative outcome)

We have already used the `Match` function from the `Matching` package to develop a matched sample. Given this, we need only run the `psens` function from the `rbounds` package to obtain sensitivity results.

```
X <- toy$linps ## matching on the linear propensity score
Tr <- as.logical(toy$treated)
Y <- toy$out1.cost
match1 <- Match(Tr=Tr, X=X, Y = Y, M = 1, replace=FALSE, ties=FALSE)
summary(match1)
```

```
Estimate... 9.7786
SE..... 1.6137
T-stat..... 6.0599
p.val..... 1.3622e-09
```

```
Original number of observations..... 400
Original number of treated obs..... 140
Matched number of observations..... 140
Matched number of observations (unweighted). 140
```

```
psens(match1, Gamma = 5, GammaInc = 0.25)
```

Rosenbaum Sensitivity Test for Wilcoxon Signed Rank P-Value

```
Unconfounded estimate .... 0
```

Gamma	Lower bound	Upper bound
1.00	0	0.0000
1.25	0	0.0000
1.50	0	0.0003
1.75	0	0.0029
2.00	0	0.0155
2.25	0	0.0510
2.50	0	0.1204
2.75	0	0.2243
3.00	0	0.3520
3.25	0	0.4867
3.50	0	0.6132
3.75	0	0.7214
4.00	0	0.8071
4.25	0	0.8711
4.50	0	0.9165
4.75	0	0.9473
5.00	0	0.9676

Note: Gamma is Odds of Differential Assignment To
Treatment Due to Unobserved Factors

If the study were free of hidden bias, that is, if $\Gamma = 1$, then there would be **strong** evidence that the treated patients had higher costs, and the specific Wilcoxon signed rank test we’re looking at here shows a p value < 0.0001 . The sensitivity analysis we’ll conduct now asks how this conclusion might be changed by hidden biases of various magnitudes, depending on the significance level we plan to use in our test.

Specifying The Threshold Γ value

From the output above, find the Γ value where the upper bound for our p value slips from “statistically significant” to “not significant” territory.

- We’re doing a two-tailed test, with a 95% confidence level, so the Γ statistic for this situation is between 2.0 and 2.25, since that is the point where the upper bound for the p value crosses the threshold of $\alpha/2 = 0.025$.

So this study's conclusion (that treated patients had significantly higher costs) would still hold even in the face of a hidden bias with $\Gamma = 2$, but not with $\Gamma = 2.25$.

The tipping point for the sensitivity parameter is a little over 2.0. To explain away the observed association between treatment and this outcome (cost), a hidden bias or unobserved covariate would need to increase the odds of treatment by more than a factor of $\Gamma = 2$.

Returning to the output:

- If instead we were doing a one-tailed test with a 90% confidence level, then the Γ statistic would be between 2.25 and 2.50, since that is where the upper bound for the p value crosses $\alpha = 0.10$.

Interpreting Γ appropriately

Γ tells you only *how big a bias is needed to change the answer*. By itself, it says NOTHING about the likelihood that a bias of that size is present in your study, except that, of course, smaller biases hide more effectively than large ones, on average.

- In some settings, we'll think of Γ in terms of small (< 1.5), modest (1.5 - 2.5), moderate (2.5 - 4) and large (> 4) hidden bias requirements. But these are completely arbitrary distinctions, and I can provide no good argument for their use.

The **only** defense against hidden bias affecting your conclusions is to try to reduce the potential for hidden bias in the first place. We work on this via careful design of observational studies, especially by including as many different dimensions of the selection problem as possible in your propensity model.

Alternative Descriptions of Γ

As we see in Chapter 9 of Rosenbaum's *Observation and Experiment*, we can describe a $\Gamma = 2$ as being equivalent to a range of potential values of Θ_p from 0.33 to 0.67, and values of $\Lambda = 3$ and $\Delta = 5$. Θ_p provides an estimate of the chance that the first person in a pair is the treated subject. Λ and Δ refer to the amplification of sensitivity analysis, with reference to a spurious association between treatment received and outcome observed in the absence of a treatment effect. The odds that the first person in a pair is treated rather than control is bounded by Λ and $1/\Lambda$. The parameter Δ defines the odds that the paired difference in outcomes is greater than 0 (as compared to less than 0) if there is in fact no treatment effect.

An Alternate Approach - the Hodges-Lehman estimate

```
hlsens(match1)
```

```
Rosenbaum Sensitivity Test for Hodges-Lehmann Point Estimate
```

```
Unconfounded estimate .... 10
```

Gamma	Lower bound	Upper bound
1	10.00000	10.0
2	4.00000	16.1
3	0.49997	19.1
4	-1.50000	21.6
5	-3.50000	23.1
6	-4.50000	24.1

```
Note: Gamma is Odds of Differential Assignment To
```

Treatment Due to Unobserved Factors

If the Γ value is 2.0, then this implies that the Hodges-Lehmann estimate might be as low as 4 or as high as 16.1 (it is 10.0 in the absence of hidden bias in this case - when $\Gamma = 0$.)

What about other types of outcomes?

The `rbounds` package can evaluate binary outcomes using the `binarysens` and `Fishersens` functions.

Survival outcomes can be assessed, too, but not, I believe, using `rbounds` unless there is no censoring. Some time back, I built a spreadsheet for this task, which I'll be happy to share.

What about when we match 1:2 or 1:3 instead of 1:1?

The `mcontrol` function in the `rbounds` package can be helpful in such a setting.

Wrapup

If you run this script, you'll wind up with a version of the `toy` tibble that contains 200 observations on 28 variables, along with a `toy.codebook` list.

You'll also have two new functions, called `szd` and `rubin3`, that, with some modification, may be useful elsewhere.

To drop everything else in the global environment created by this Markdown file, run the code that follows.

```
rm(list = c("adj.m.out1", "adj.m.out1.tidy", "adj.m.out2", "adj.m.out2_tidy",
"adj.m.out3", "adj.m.out3_tidy", "adj.reg.out1", "adj.reg.out2",
"adj.reg.out3", "adj.s.out3", "adj_out1", "adj_out2", "adj_out3",
"adjout1.wt1", "adjout1.wt2", "adjout1.wt3", "adjout2.wt1", "adjout2.wt2",
"adjout2.wt3", "adjout3.wt1", "adjout3.wt2", "adjout3.wt3", "alert",
"b", "bal.after.wts1", "bal.after.wts2", "bal.before.wts1", "bal.before.wts2",
"bal.wts1", "bal.wts2", "balance.ate.weights", "balance.att.weights",
"cov.sub", "covlist", "covnames", "covs", "d.all", "d.q1", "d.q2",
"d.q3", "d.q4", "d.q5", "decim", "dr.out1.wt1", "dr.out1.wt2",
"dr.out1.wt3", "dr.out2.wt1", "dr.out2.wt2", "dr.out2.wt3", "dr.out3.wt1",
"dr.out3.wt2", "dr.out3.wt3", "dr_ate_out1", "dr_ate_out2", "dr_ate_out3",
"dr_att_out1", "dr_att_out2", "dr_att_out3", "dr_twangatt_out1",
"dr_twangatt_out2", "dr_twangatt_out3", "est.st", "factorlist",
"i", "match_szd", "match_vrat", "match1", "match1.out1", "match1.out1.ATE",
"match1_out2", "matched_mixedmodel.out1", "matches", "mb1", "p",
"post.szd", "post.vratio", "pre.szd", "pre.vratio", "ps.toy",
"psmodel", "quin1", "quin1.out1", "quin1.out2", "quin2", "quin2.out1",
"quin2.out2", "quin3", "quin3.out1", "quin3.out2", "quin4", "quin4.out1",
"quin4.out2", "quin5", "quin5.out1", "quin5.out2", "res_matched_1",
"res_unadj_1", "res_unadj_2_oddsratio", "res_unadj_2_or", "res_unadj_2_riskdiff",
"res_unadj_3", "rubin1.match", "rubin1.q1", "rubin1.q2", "rubin1.q3",
"rubin1.q4", "rubin1.q5", "rubin1.sub", "rubin1.unadj", "rubin2.match",
"rubin2.q1", "rubin2.q2", "rubin2.q3", "rubin2.q4", "rubin2.q5",
"rubin2.sub", "rubin2.unadj", "rubin3.both", "rubin3.matched",
"rubin3.q1", "rubin3.q2", "rubin3.q3", "rubin3.q4", "rubin3.q5",
"rubin3.unadj", "se.q1", "se.q2", "se.q3", "se.q4", "se.q5",
```



```
"se.st", "strat.result1", "strat.result2", "strat.result3",  
"temp", "toy.matchedsample", "toy.rubin3",  
"toy.szd", "toy_df", "toywt1.design", "toywt2.design", "toywt3.design",  
"Tr", "unadj.out1", "unadj.out2", "unadj.out3", "varlist", "wt_ate_results1",  
"wt_ate_results2", "wt_ate_results3", "wt_att_results1", "wt_att_results2",  
"wt_att_results3", "wt_twangatt_results1", "wt_twangatt_results2",  
"wt_twangatt_results3", "wts3", "X", "Y"))
```